



Parallel Algorithms

Patrick Cozzi

University of Pennsylvania

CIS 565 - Fall 2015

Agenda

■ Parallel Algorithms

- Parallel Reduction
- Scan (Naive and Work-Efficient)
- Stream Compression
- Summed *A*rea *T*ables
- Radix Sort

Parallel Reduction

- Given an array of numbers, design a parallel algorithm to find the sum.
- Consider:
 - *Arithmetic intensity*: compute to memory access ratio



Parallel Reduction

- Given an array of numbers, design a parallel algorithm to find:
 - The sum
 - The maximum value
 - The product of values
 - The average value
- How different are these algorithms?

Parallel Reduction

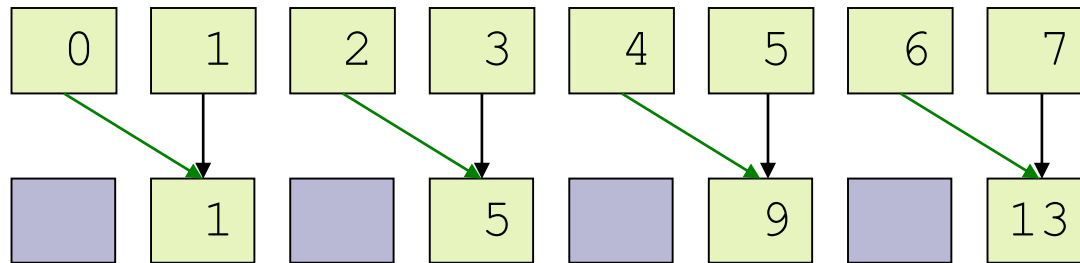
- *Reduction*: An operation that computes a single result from a set of data
- *Parallel Reduction*: Do it in parallel.
Obviously

Parallel Reduction

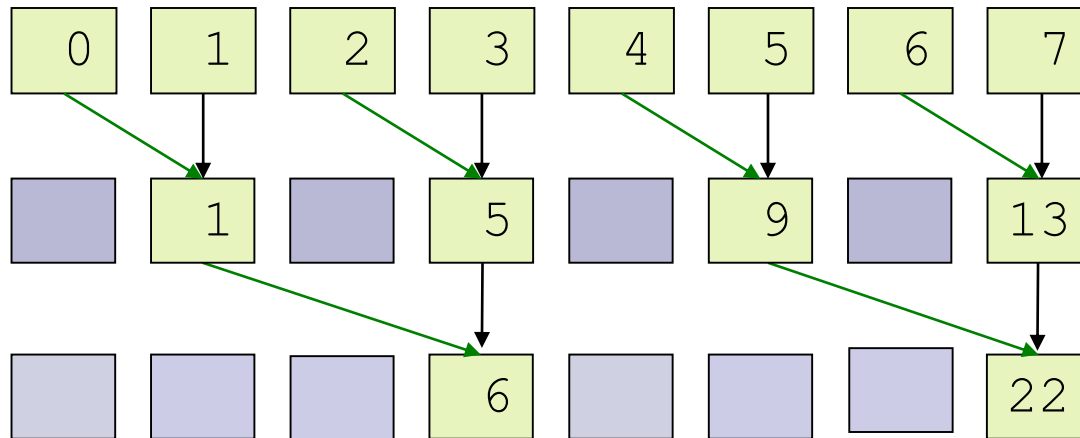
- Example. Find the sum:

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

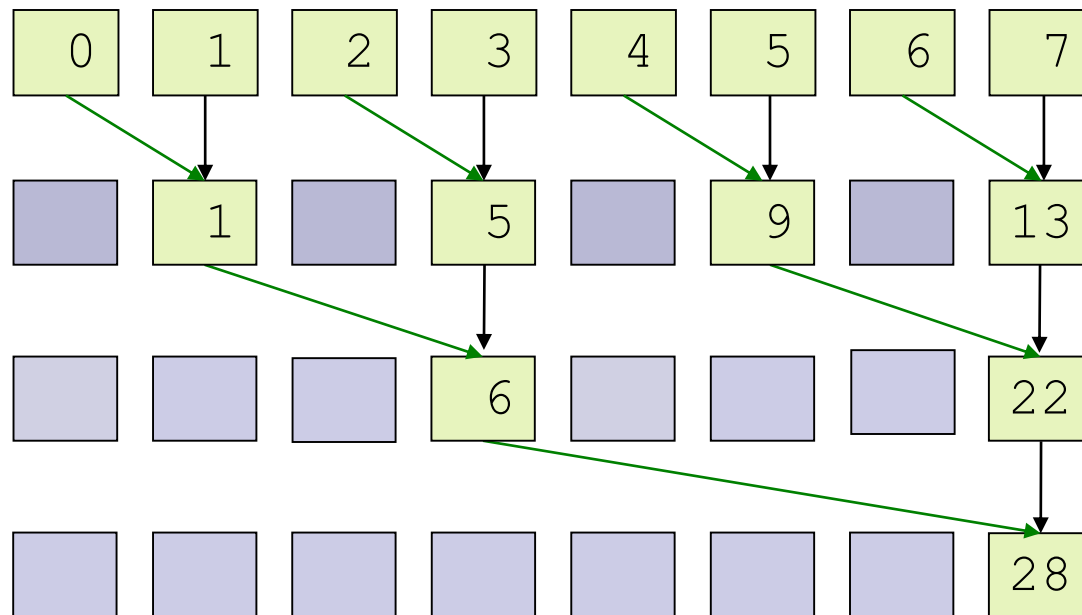
Parallel Reduction



Parallel Reduction

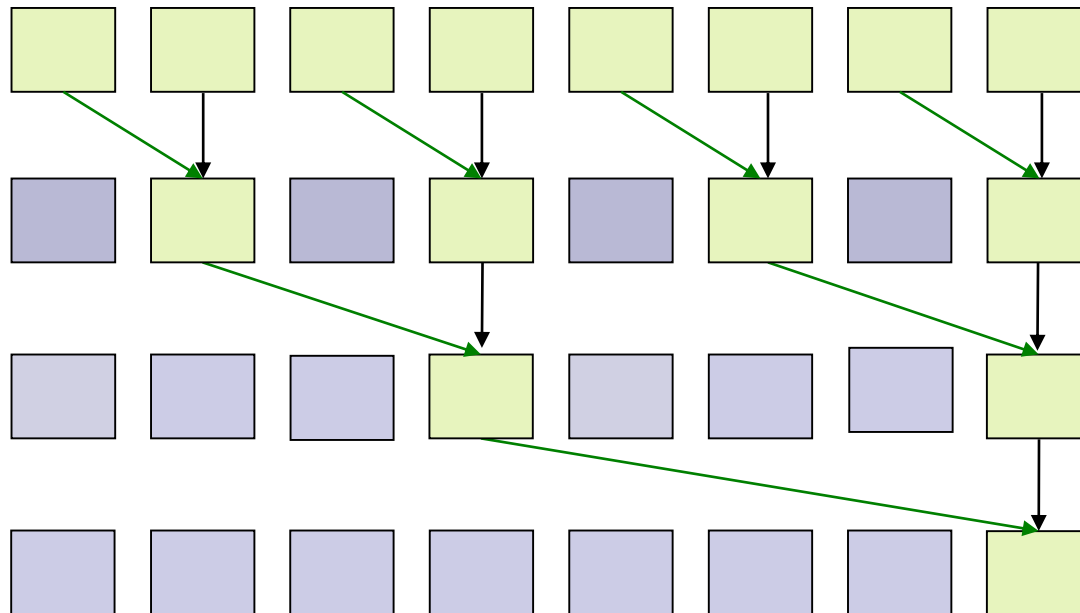


Parallel Reduction



Parallel Reduction

- Similar to brackets for a basketball tournament
- $\log(n)$ passes for n elements



Parallel Reduction

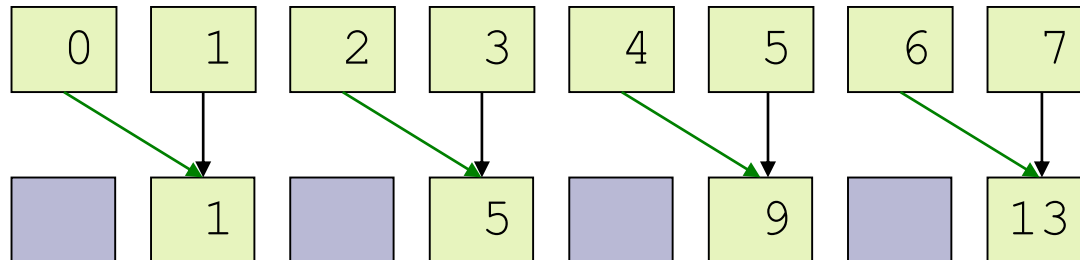
■ $d = 0, 2^{d+1} = 2$

■ $2^{d+1} - 1 = 1$

■ $2^d - 1 = 0$

```
for d = 0 to log2n - 1
  for all k = 0 to n - 1 by 2d+1 in parallel
    x[k + 2d+1 - 1] += x[k + 2d - 1];
```

```
// In this pass, for k = (0, 2, 4, 6)
//   x[k + 1] += x[k];
```



Parallel Reduction

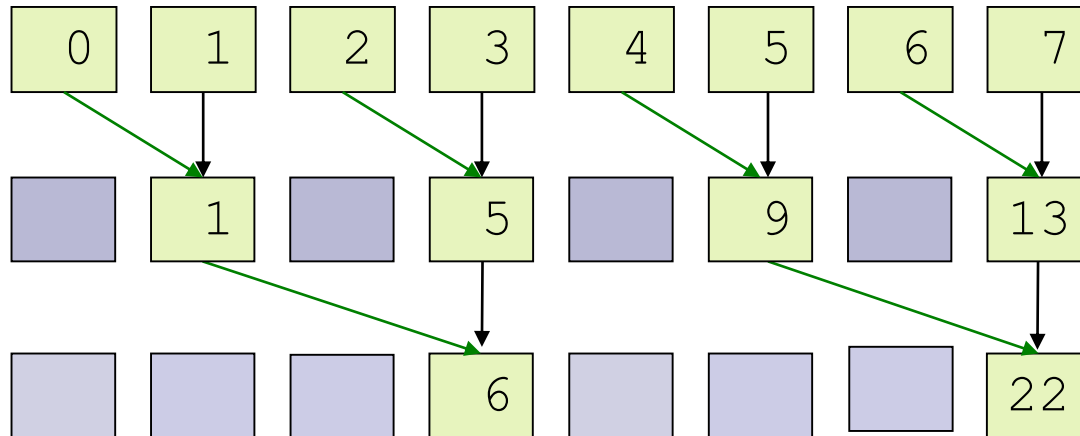
■ $d = 1, 2^{d+1} = 4$

■ $2^{d+1} - 1 = 3$

■ $2^d - 1 = 1$

```
for d = 0 to log2n - 1
  for all k = 0 to n - 1 by 2d+1 in parallel
    x[k + 2d+1 - 1] += x[k + 2d - 1];
```

```
// In this pass, for k = (0, 4)
//   x[k + 3] += x[k + 1];
```



Parallel Reduction

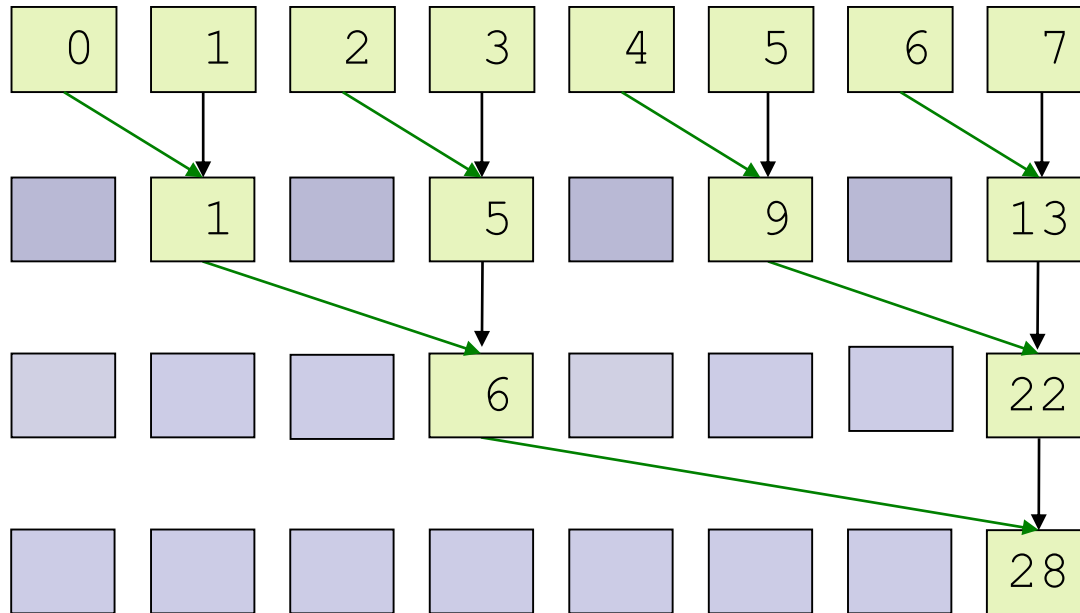
■ $d = 2, 2^{d+1} = 8$

■ $2^{d+1} - 1 = 7$

■ $2^d - 1 = 3$

```
for d = 0 to  $\log_2 n - 1$ 
  for all k = 0 to n - 1 by  $2^{d+1}$  in parallel
     $x[k + 2^{d+1} - 1] += x[k + 2^d - 1];$ 
```

```
// In this pass, for k = (0)
//    $x[k + 7] += x[k + 3];$ 
```

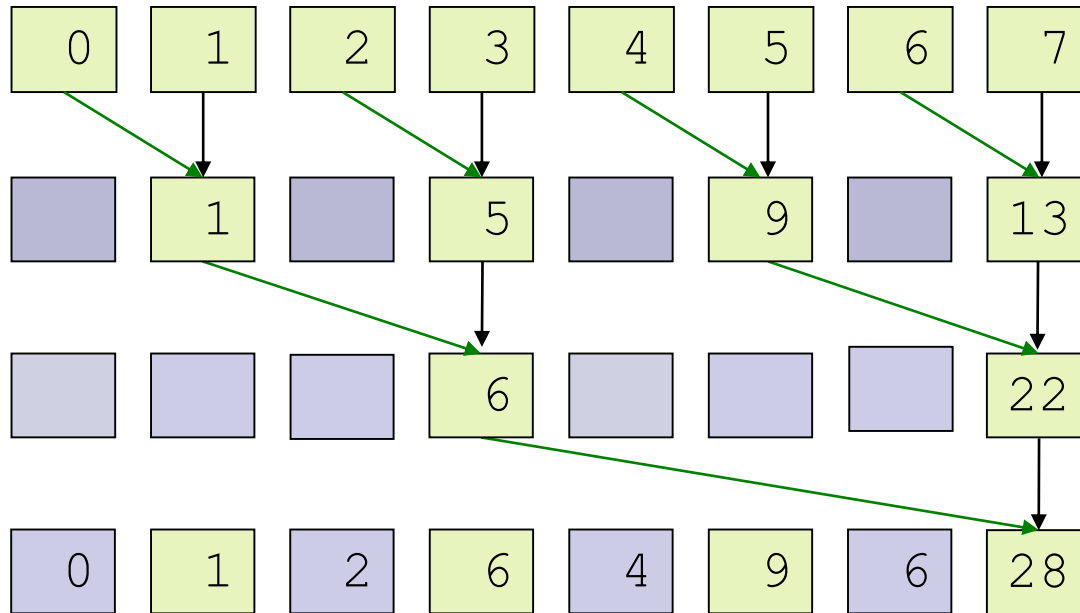


Parallel Reduction

- Note the $+=$

```
for d = 0 to  $\log_2 n - 1$   
  for all k = 0 to n - 1 by  $2^{d+1}$  in parallel  
     $x[k + 2^{d+1} - 1] += x[k + 2^d - 1];$ 
```

- The array is modified in place



All-Prefix-Sums

■ *All-Prefix-Sums*

□ Input

- Array of n elements: $[a_0, a_1, \dots, a_{n-1}]$
- Binary associate operator: \oplus
- Identity: I

□ Outputs the array: $[I, a_0, (a_0 \oplus a_1), \dots, (a_0 \oplus a_1 \oplus \dots \oplus a_{n-2})]$

All-Prefix-Sums

■ Example

□ If \oplus is addition, the array

■ [3 1 7 0 4 1 6 3]

□ is transformed to

■ [0 3 4 11 11 15 16 22]

■ Seems sequential, but there is an efficient parallel solution

Scan

- **Exclusive Scan**: Element j of the result does not include element j of the input:

- In: [3 1 7 0 4 1 6 3]

- Out: [0 3 4 11 11 15 16 22]

- **Inclusive Scan (Prescan)**: All elements including j are summed

- In: [3 1 7 0 4 1 6 3]

- Out: [3 4 11 11 15 16 22 25]

Scan

- How do you generate an *exclusive scan* from an *inclusive scan*?

- Input: [3 1 7 0 4 1 6 3]
- Inclusive: [3 4 11 11 15 16 22 25]
- Exclusive: [0 3 4 11 11 15 16 22]
 - // Shift right, insert identity

- How do you go in the opposite direction?

Scan

■ Use cases

- *Stream compaction*
- *Summed-area tables* for variable width image processing
- *Radix sort*
- ...

Scan

- Used to convert certain sequential computation into equivalent parallel computation

Sequential	Parallel
<pre>01. out[0] = 0; 02. for j from 1 to n do 03. out[j] = out[j-1] + f(in[j-1]);</pre>	<pre>01. forall j in parallel do 02. temp[j] = f(in[j]); 03. all_prefix_sums(out, temp);</pre>

Scan

- Design a parallel algorithm for **inclusive** scan

- In: [3 1 7 0 4 1 6 3]

- Out: [3 4 11 11 15 16 22 25]

- **Consider:**

- **Total number of additions**

Scan

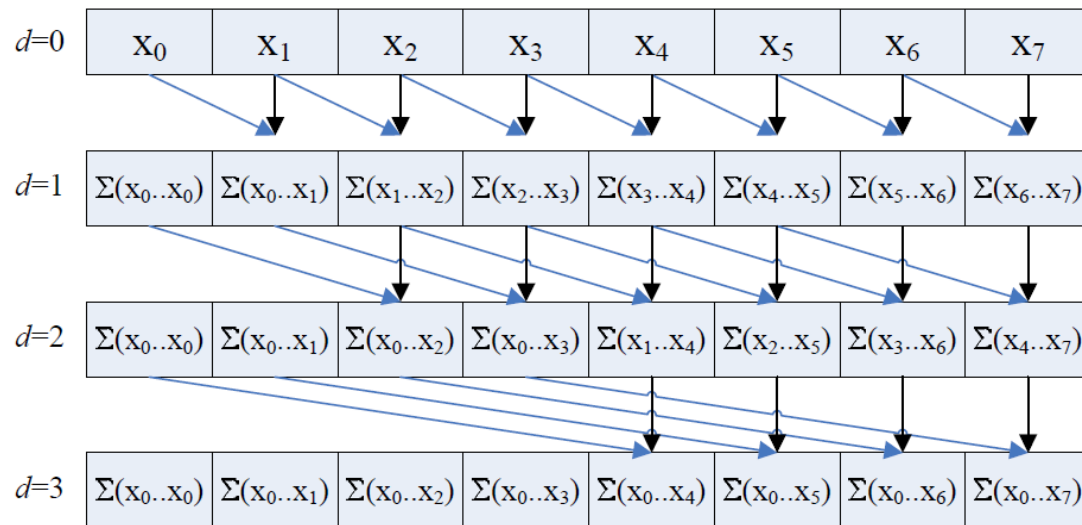
- Single thread (*Sequential Scan*) is trivial:

```
01. out[0] := 0
02. for k := 1 to n do
03.     out[k] := in[k-1] + out[k-1]
```

- n adds for an array of length n
- How many adds will our parallel version have?

Scan

■ *Naive Parallel Scan*



```
for d = 1 to  $\log_2 n$ 
  for all k in parallel
    if ( $k \geq 2^{d-1}$ )
       $x[k] = x[k - 2^{d-1}] + x[k];$ 
```

- Is this exclusive or inclusive?
- Each thread
 - Writes one sum
 - Reads two values

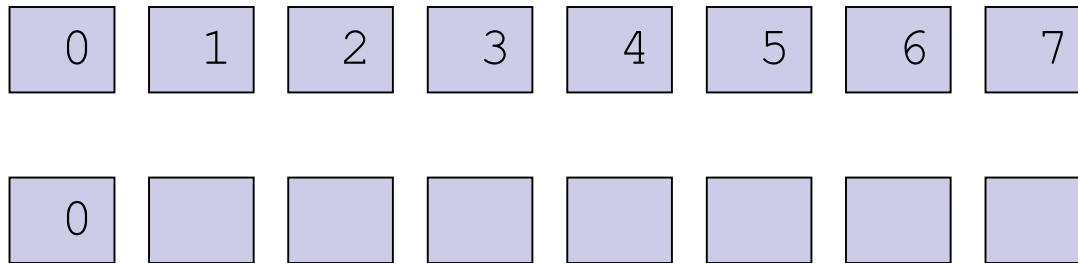
Scan

- *Naive Parallel Scan*: Input



Scan

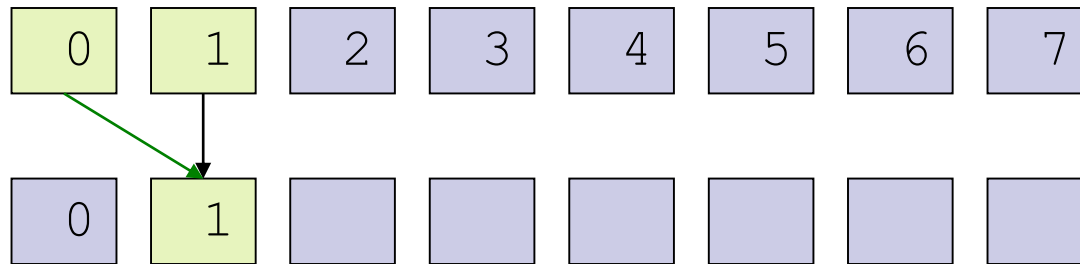
- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

Scan

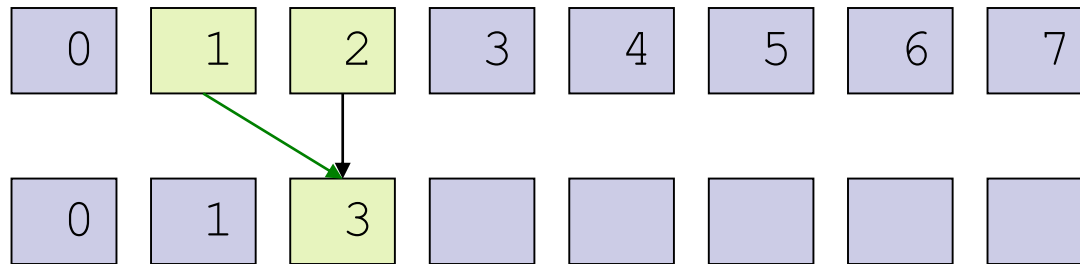
- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$



```
for d = 1 to  $\log_2 n$ 
  for all k in parallel
    if  $(k \geq 2^{d-1})$ 
       $x[k] = x[k - 2^{d-1}] + x[k];$ 
```

Scan

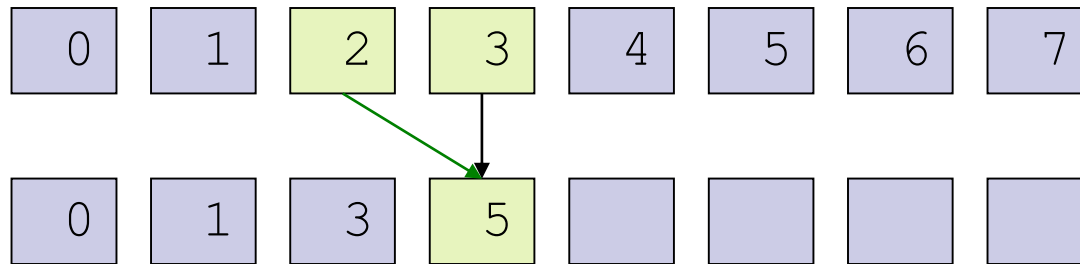
- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

Scan

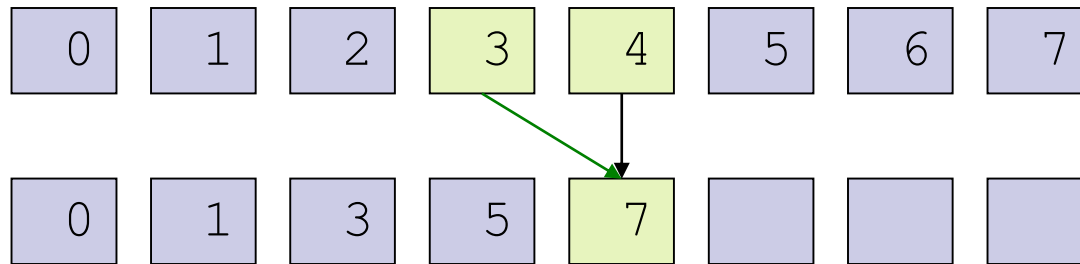
- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

Scan

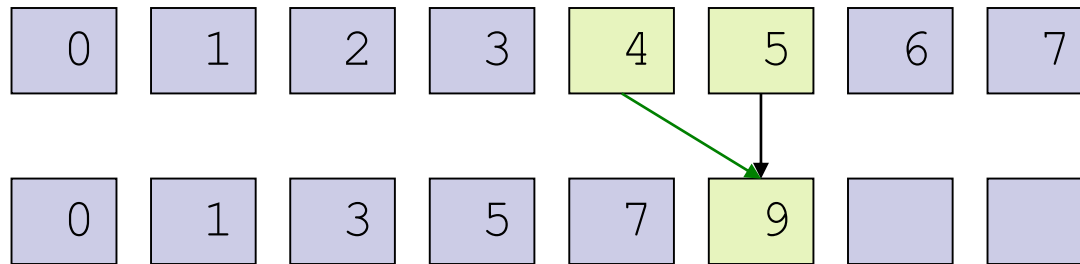
- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

Scan

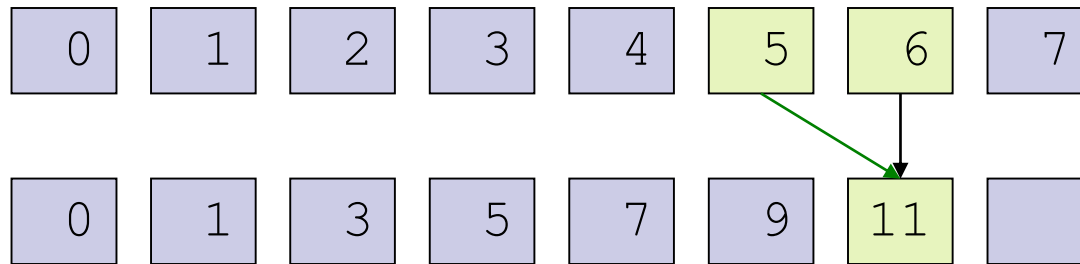
- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$



```
for d = 1 to  $\log_2 n$ 
  for all k in parallel
    if  $(k \geq 2^{d-1})$ 
      30
       $x[k] = x[k - 2^{d-1}] + x[k];$ 
```

Scan

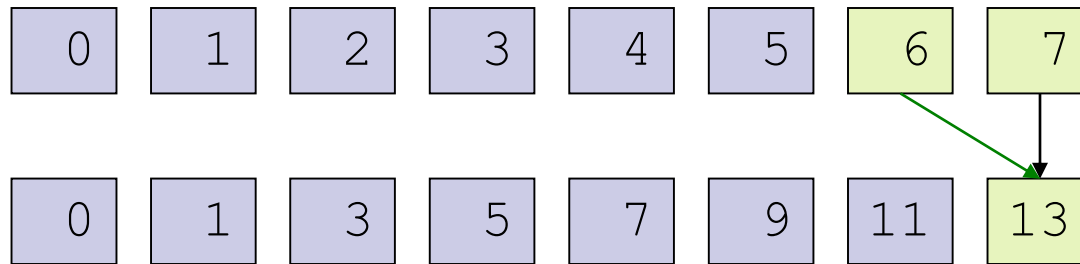
- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

Scan

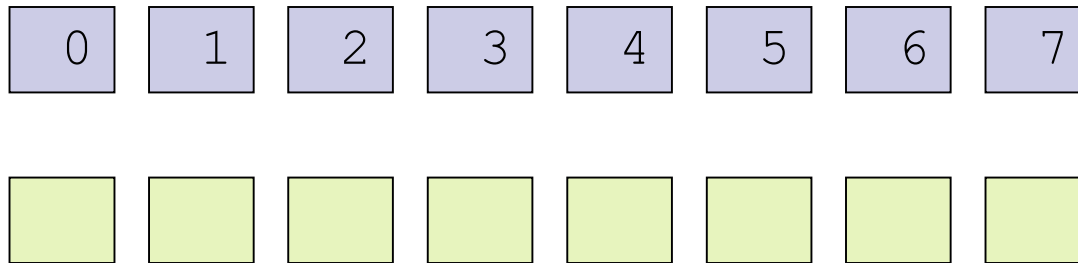
- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```


Scan

- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$

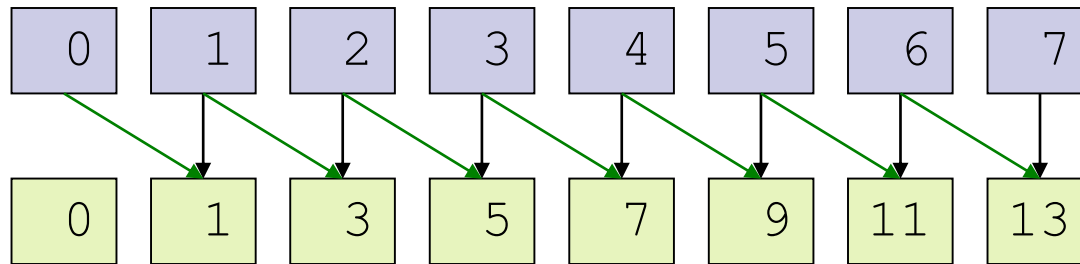


- Recall, it runs in parallel!

```
for d = 1 to  $\log_2 n$ 
  for all k in parallel
    if  $(k \geq 2^{d-1})$ 
       $x[k] = x[k - 2^{d-1}] + x[k];$ 
```

Scan

- *Naive Parallel Scan*: $d = 1, 2^{d-1} = 1$

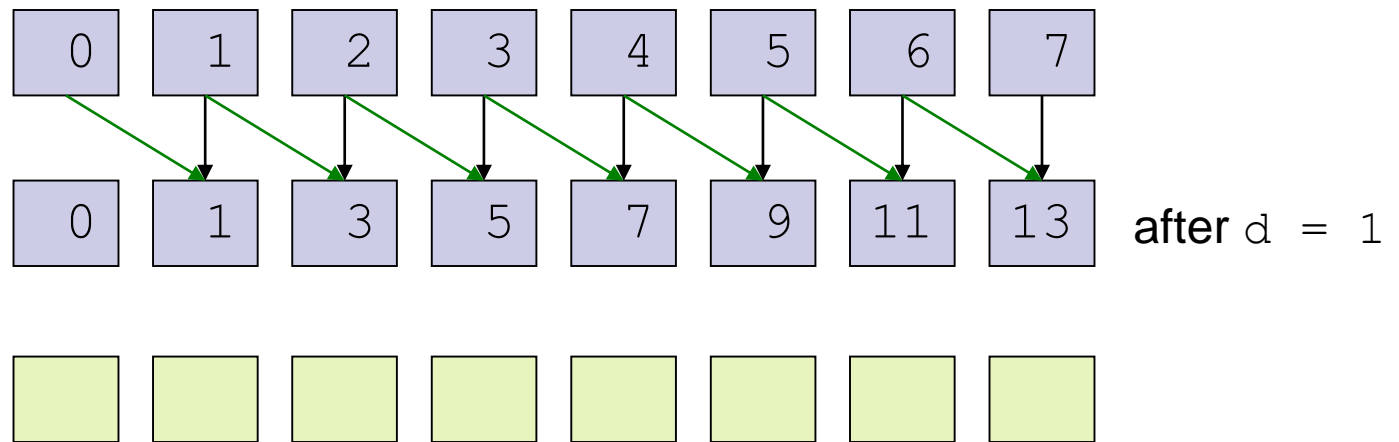


- Recall, it runs in parallel!

```
for d = 1 to  $\log_2 n$ 
  for all k in parallel
    if ( $k \geq 2^{d-1}$ )
       $x[k] = x[k - 2^{d-1}] + x[k];$ 
```

Scan

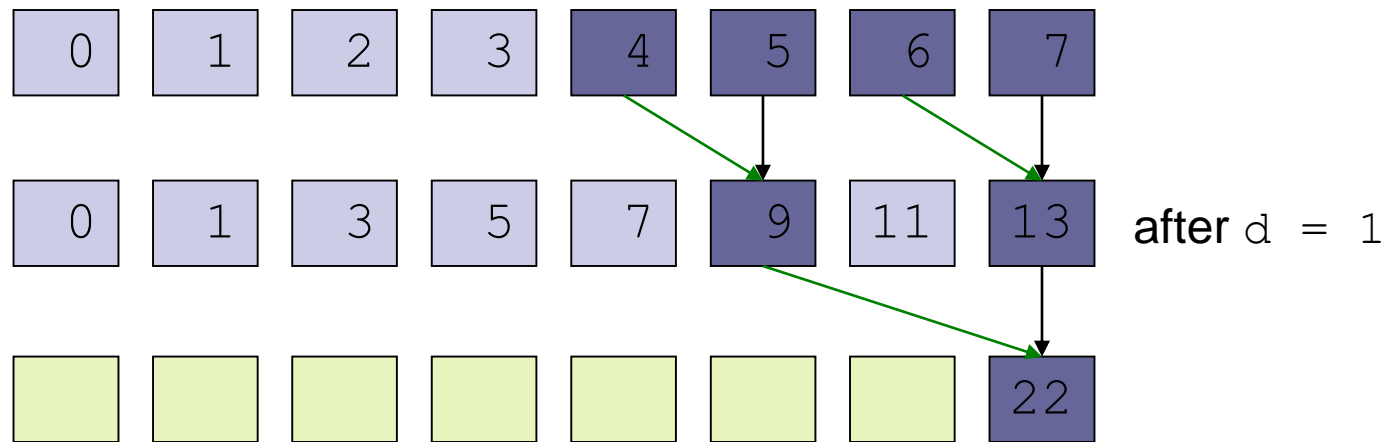
- *Naive Parallel Scan*: $d = 2, 2^{d-1} = 2$



```
for d = 1 to  $\log_2 n$ 
  for all k in parallel
    if  $(k \geq 2^{d-1})$ 
       $x[k] = x[k - 2^{d-1}] + x[k];$ 
```

Scan

- *Naive Parallel Scan*: $d = 2, 2^{d-1} = 2$

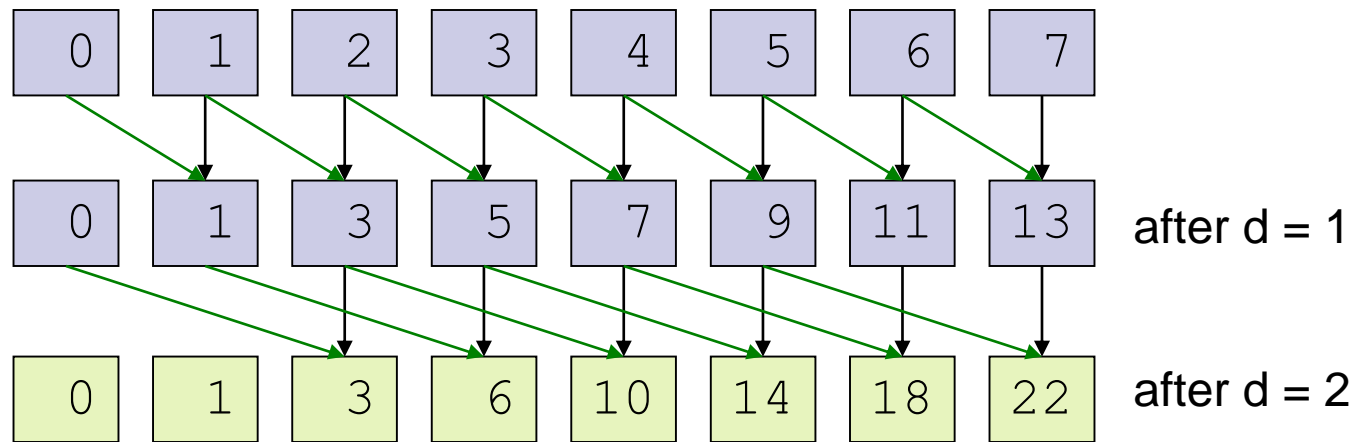


- Consider only $k = 7$

```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

Scan

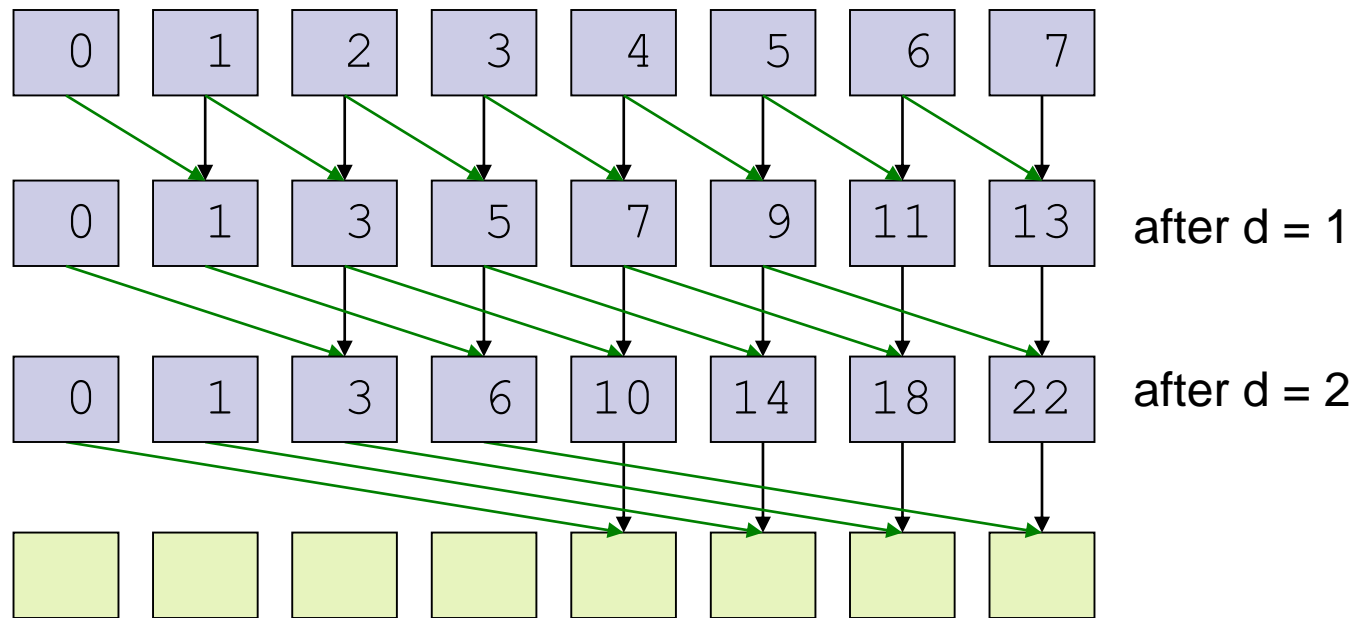
- *Naive Parallel Scan*: $d = 2, 2^{d-1} = 2$



```
for d = 1 to log2n
  for all k in parallel
    if (k >= 2d-1)
      x[k] = x[k - 2d-1] + x[k];
```

Scan

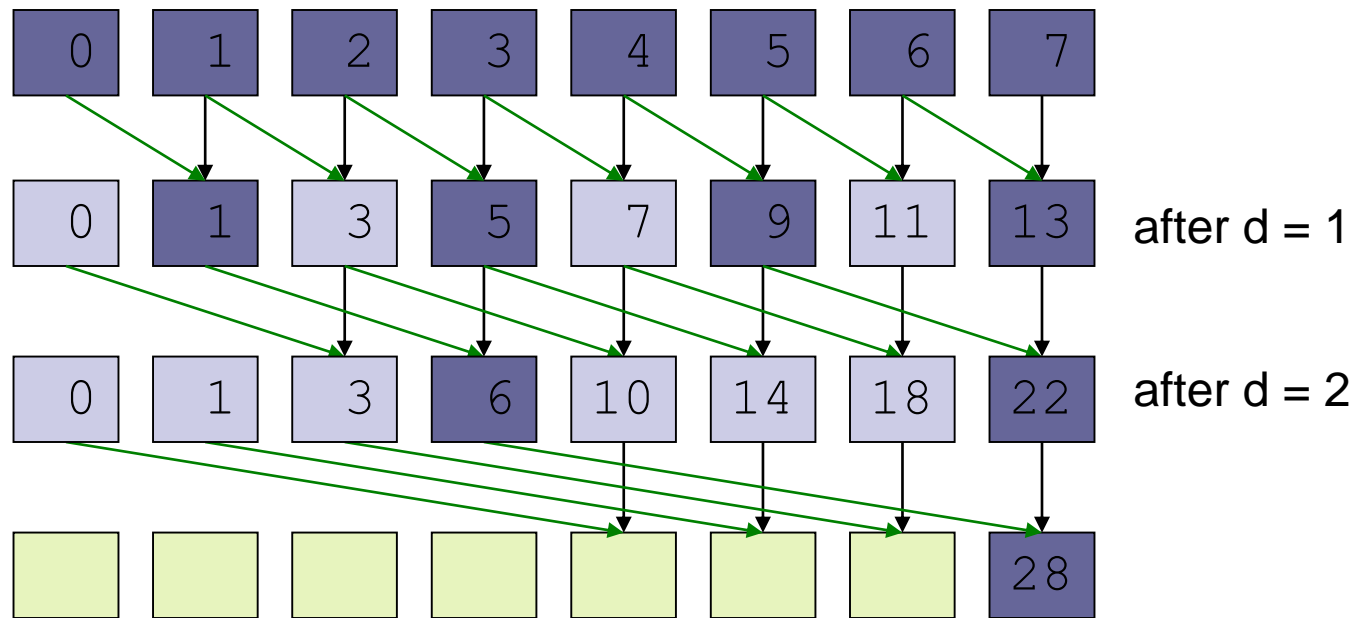
- *Naive Parallel Scan*: $d = 3$, $2^{d-1} = 4$



```
for d = 1 to  $\log_2 n$ 
  for all k in parallel
    if  $(k \geq 2^{d-1})$ 
       $x[k] = x[k - 2^{d-1}] + x[k];$ 
```

Scan

- *Naive Parallel Scan*: $d = 3, 2^{d-1} = 4$



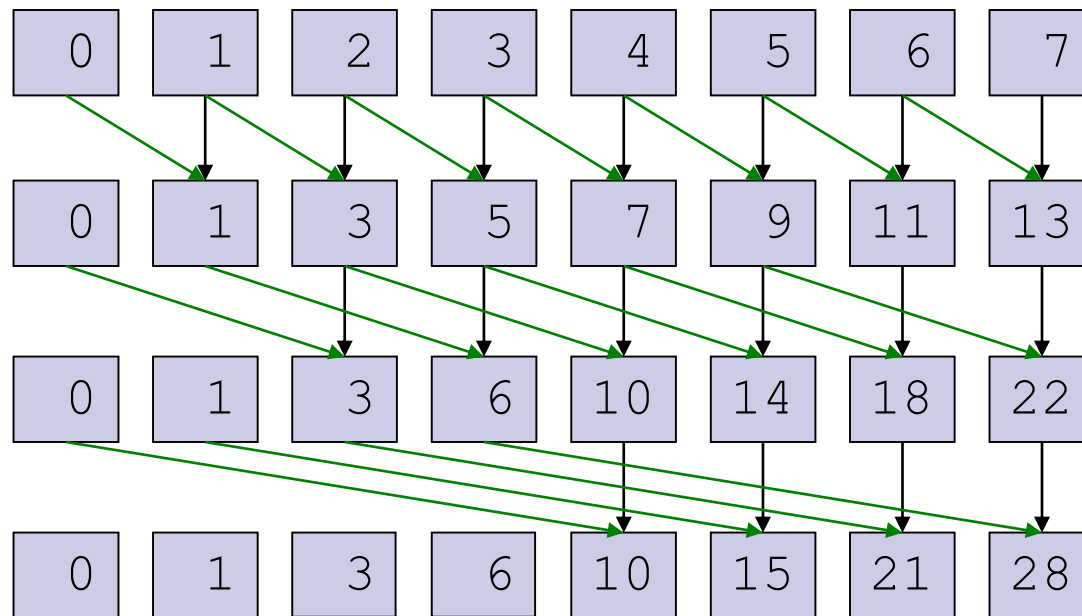
- Consider only $k = 7$

```

for d = 1 to  $\log_2 n$ 
  for all k in parallel
    if ( $k \geq 2^{d-1}$ )
       $x[k] = x[k - 2^{d-1}] + x[k];$ 
  
```

Scan

■ *Naive Parallel Scan*: Final



Work-Efficient Parallel Scan

- Number of adds

- *Sequential Scan*: $O(n)$

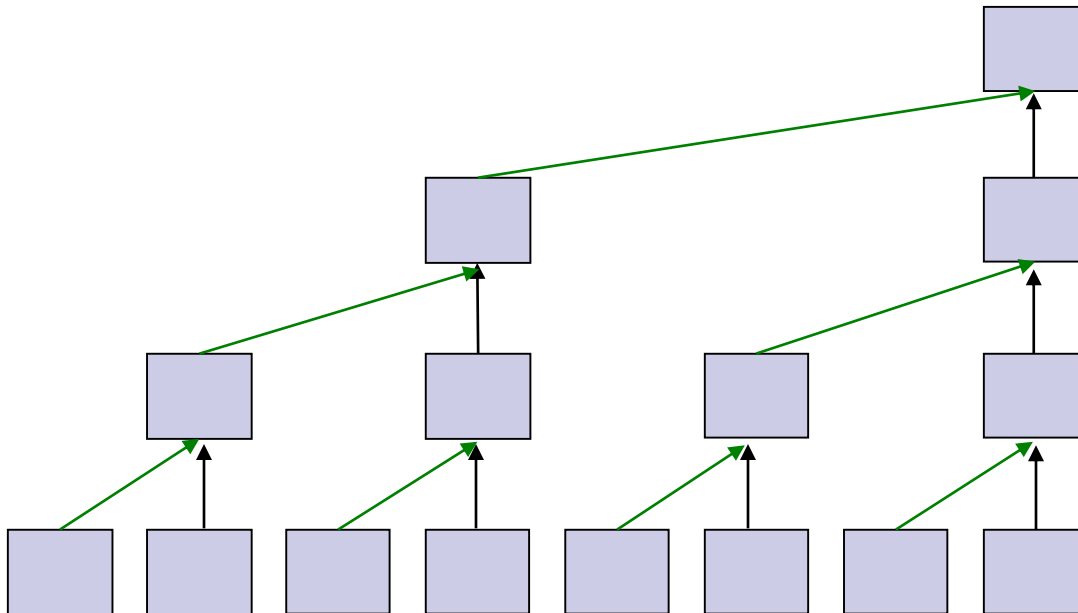
- *Naive Parallel Scan*: $O(n \log_2(n))$

- How can we make it faster?

Work-Efficient Parallel Scan

■ *Balanced binary tree*

- n leaves = $\log_2 n$ levels
- Each level, d , has 2^d nodes

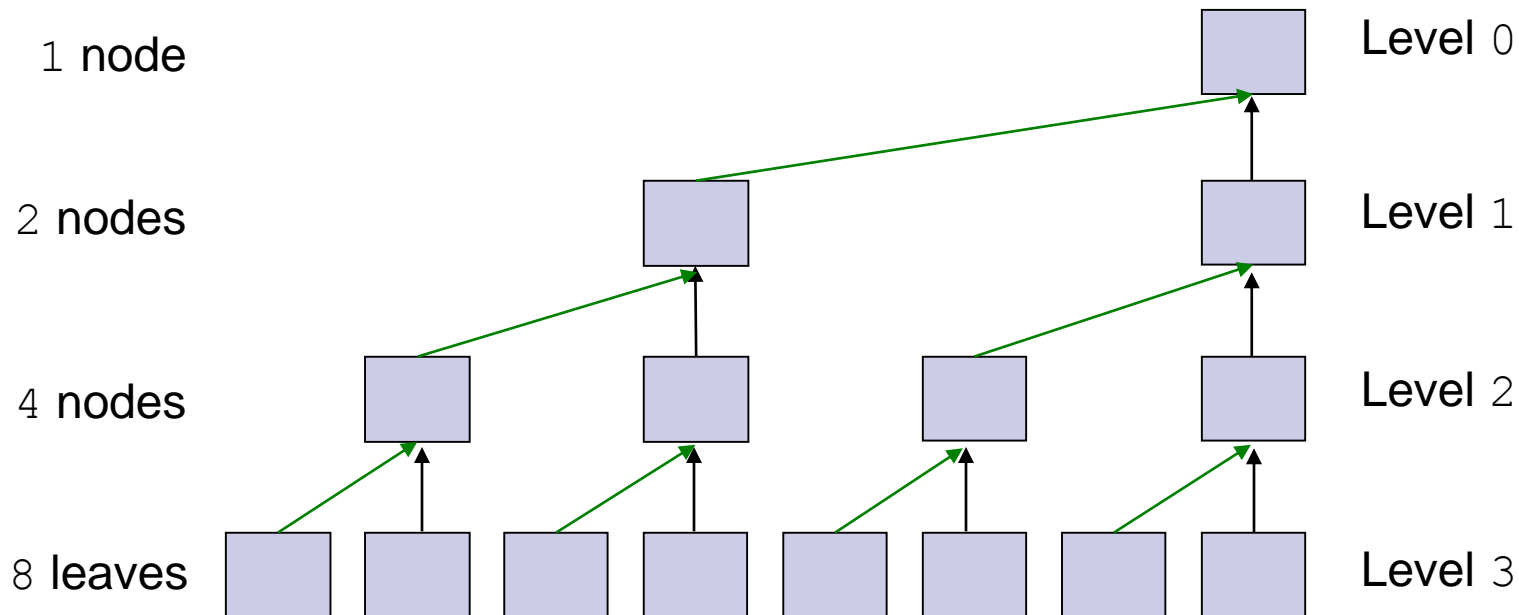


Work-Efficient Parallel Scan

■ *Balanced binary tree*

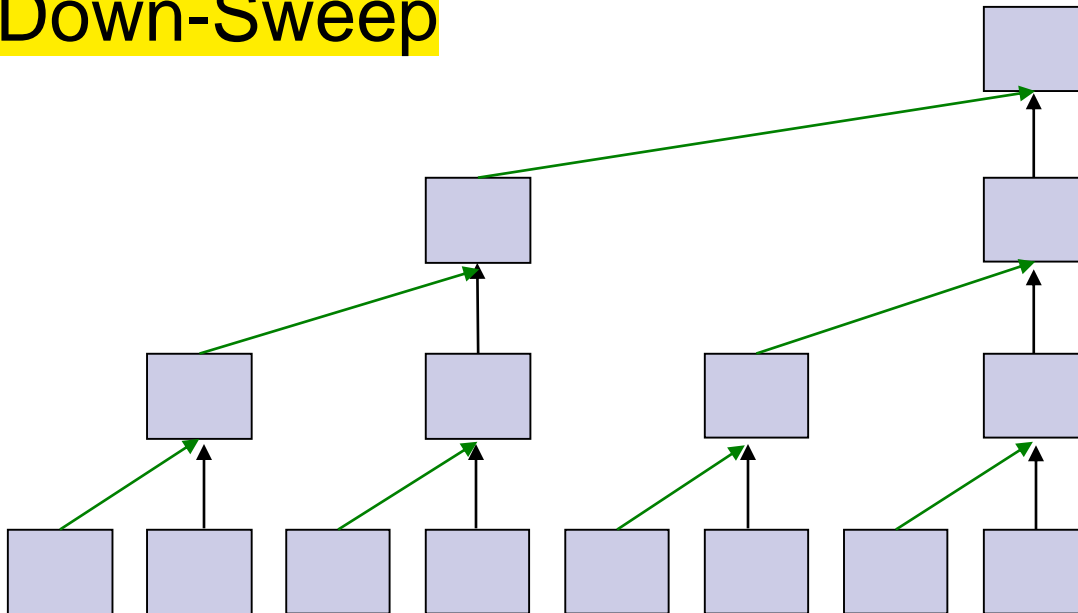
□ n leaves = $\log_2 n$ levels

□ Each level, d , has 2^d nodes



Work-Efficient Parallel Scan

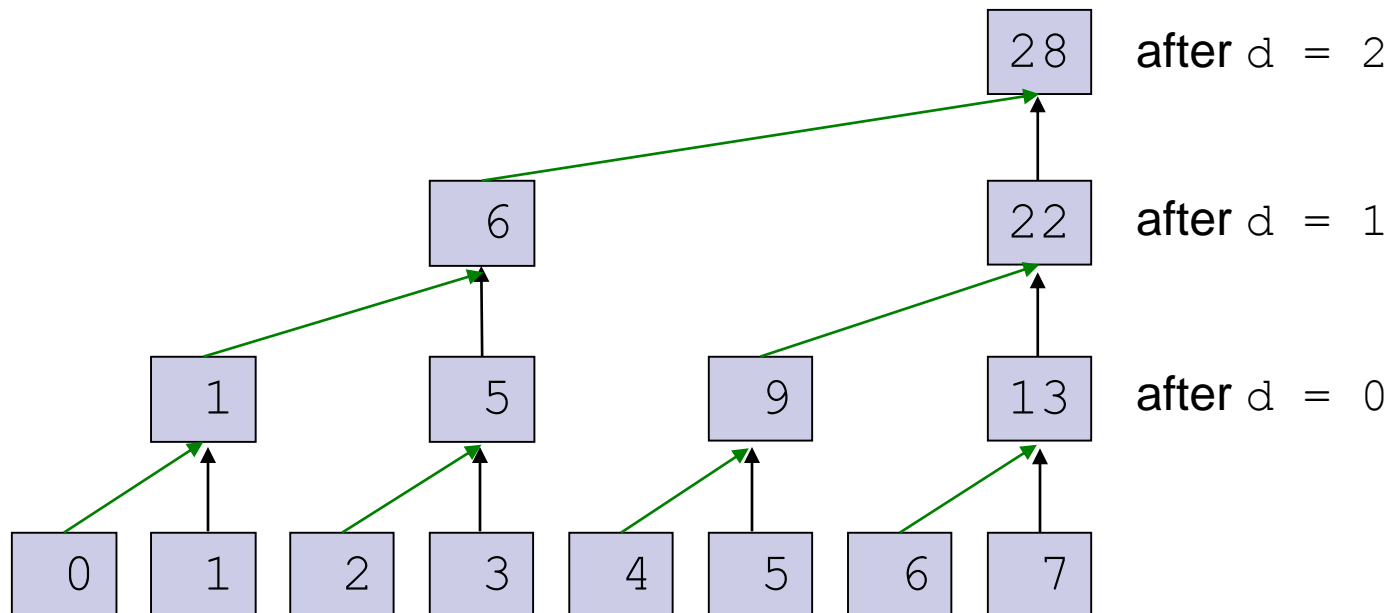
- Use a *balanced binary tree* (in concept) to perform Scan in two phases:
 - Up-Sweep (Parallel Reduction)
 - Down-Sweep



Work-Efficient Parallel Scan

■ Up-Sweep

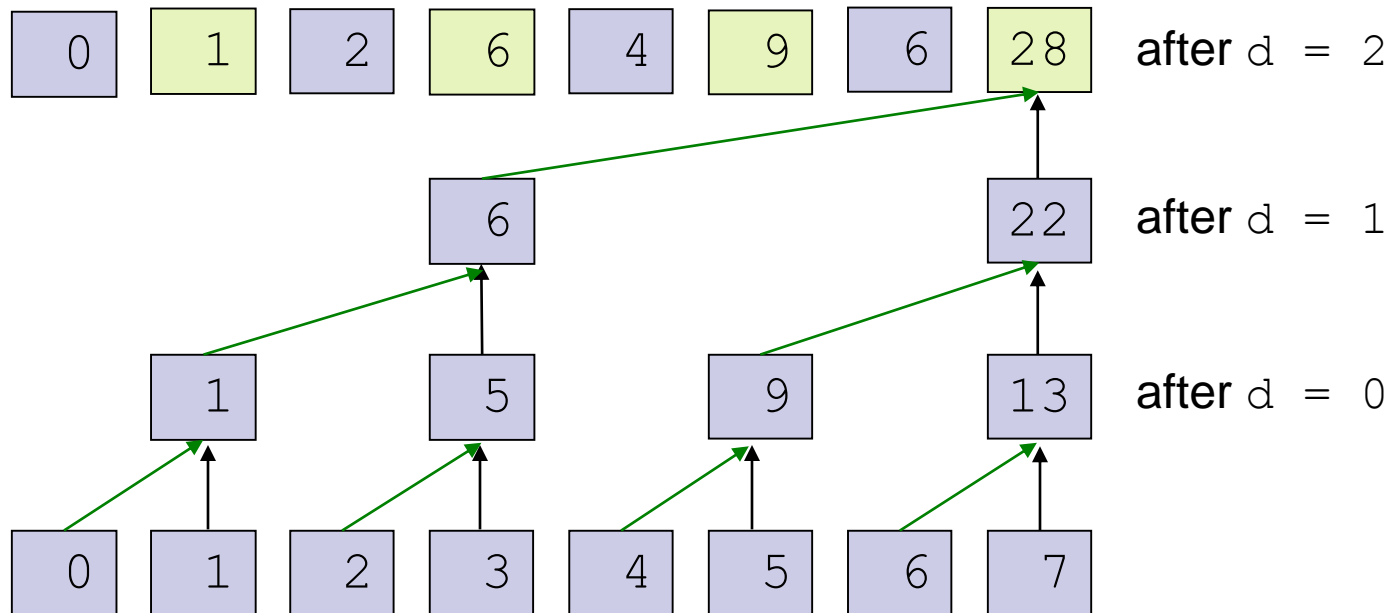
```
// Same code as our Parallel Reduction
for d = 0 to  $\log_2 n - 1$ 
  for all  $k = 0$  to  $n - 1$  by  $2^{d+1}$  in parallel
     $x[k + 2^{d+1} - 1] += x[k + 2^d - 1];$ 
```



Work-Efficient Parallel Scan

■ Up-Sweep

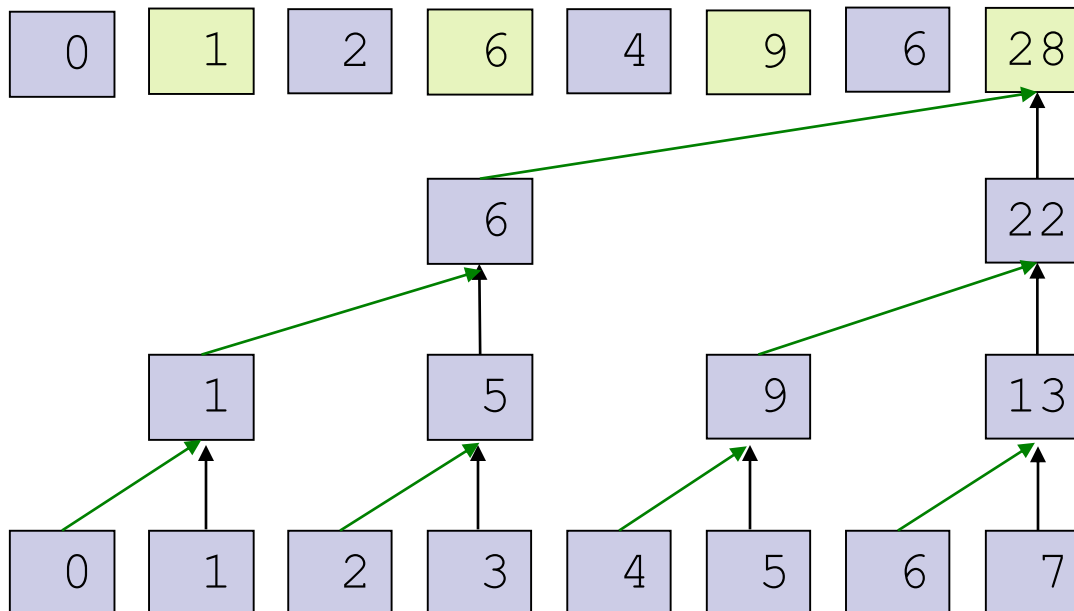
```
// Same code as our Parallel Reduction
for d = 0 to  $\log_2 n - 1$ 
  for all  $k = 0$  to  $n - 1$  by  $2^{d+1}$  in parallel
     $x[k + 2^{d+1} - 1] += x[k + 2^d - 1];$ 
```



Work-Efficient Parallel Scan

■ Down-Sweep

- “Traverse” back down tree using partial sums to build the scan in place.
 - Set **root to zero**
 - At each pass, a node passes **its value** to its **left child**, and sets the **right child** to the **sum** of the **previous left child's** value and **its value**



Work-Efficient Parallel Scan

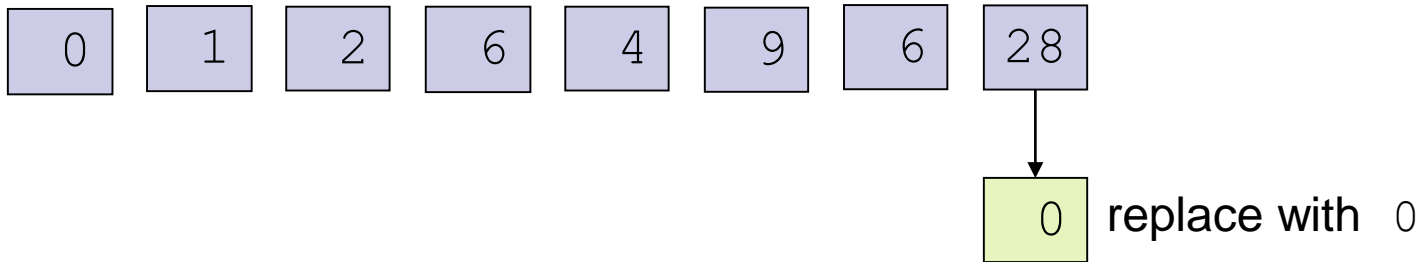
■ Down-Sweep

- “Traverse” back down tree using partial sums to build the scan in place.
 - Set root to zero
 - At each pass, a node passes its value to its left child, and sets the right child to the sum of the previous left child’s value and its value

```
x[n - 1] = 0
for d = log2n - 1 to 0
  for all k = 0 to n - 1 by 2d+1 in parallel
    t = x[k + 2d - 1];           // Save left child
    x[k + 2d - 1] = x[k + 2d+1 - 1]; // Set left child to this node's value
    x[k + 2d+1 - 1] += t;         // Set right child to old left value +
                                   // this node's value
```

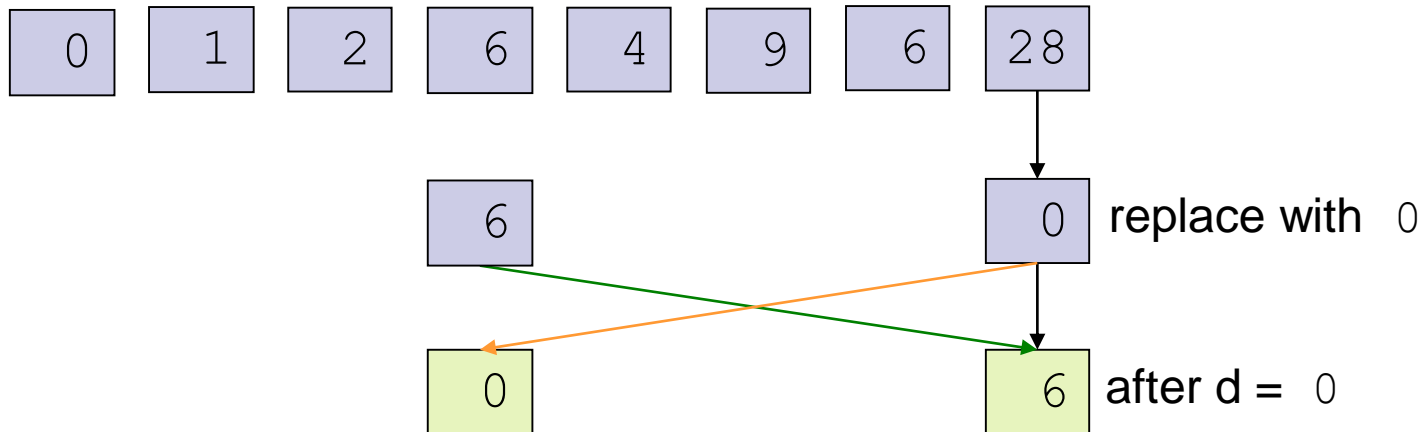

Work-Efficient Parallel Scan

■ Down-Sweep



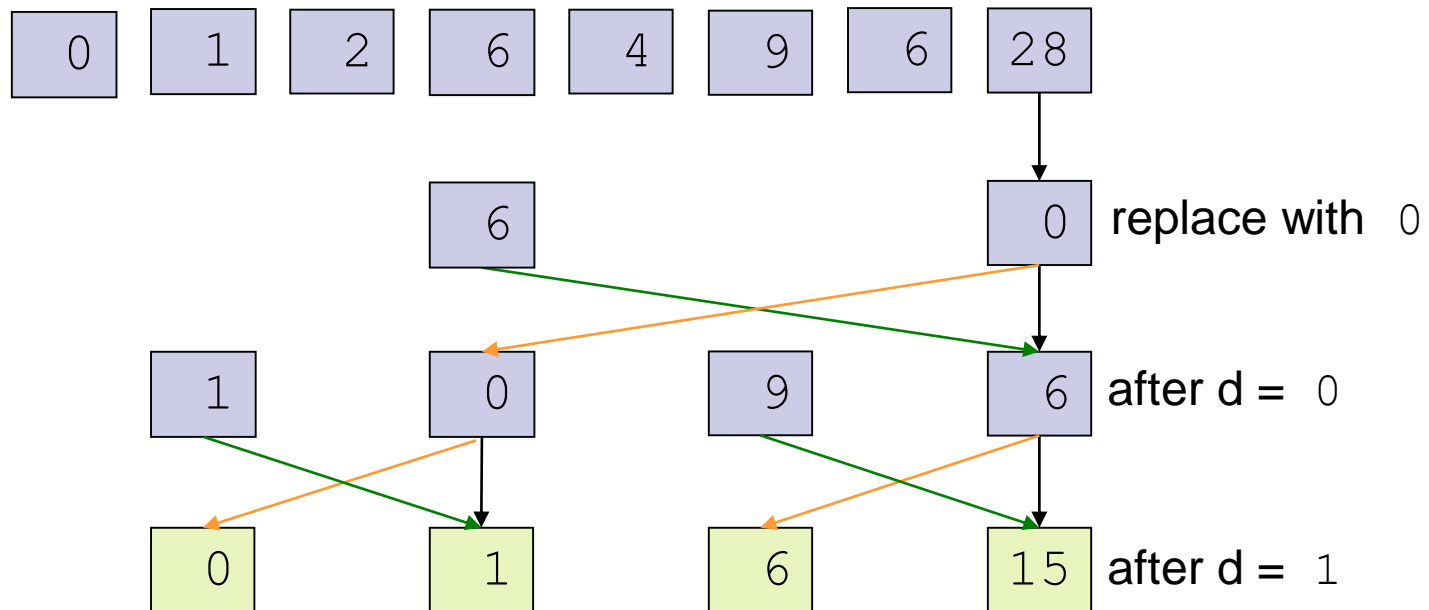
Work-Efficient Parallel Scan

■ Down-Sweep



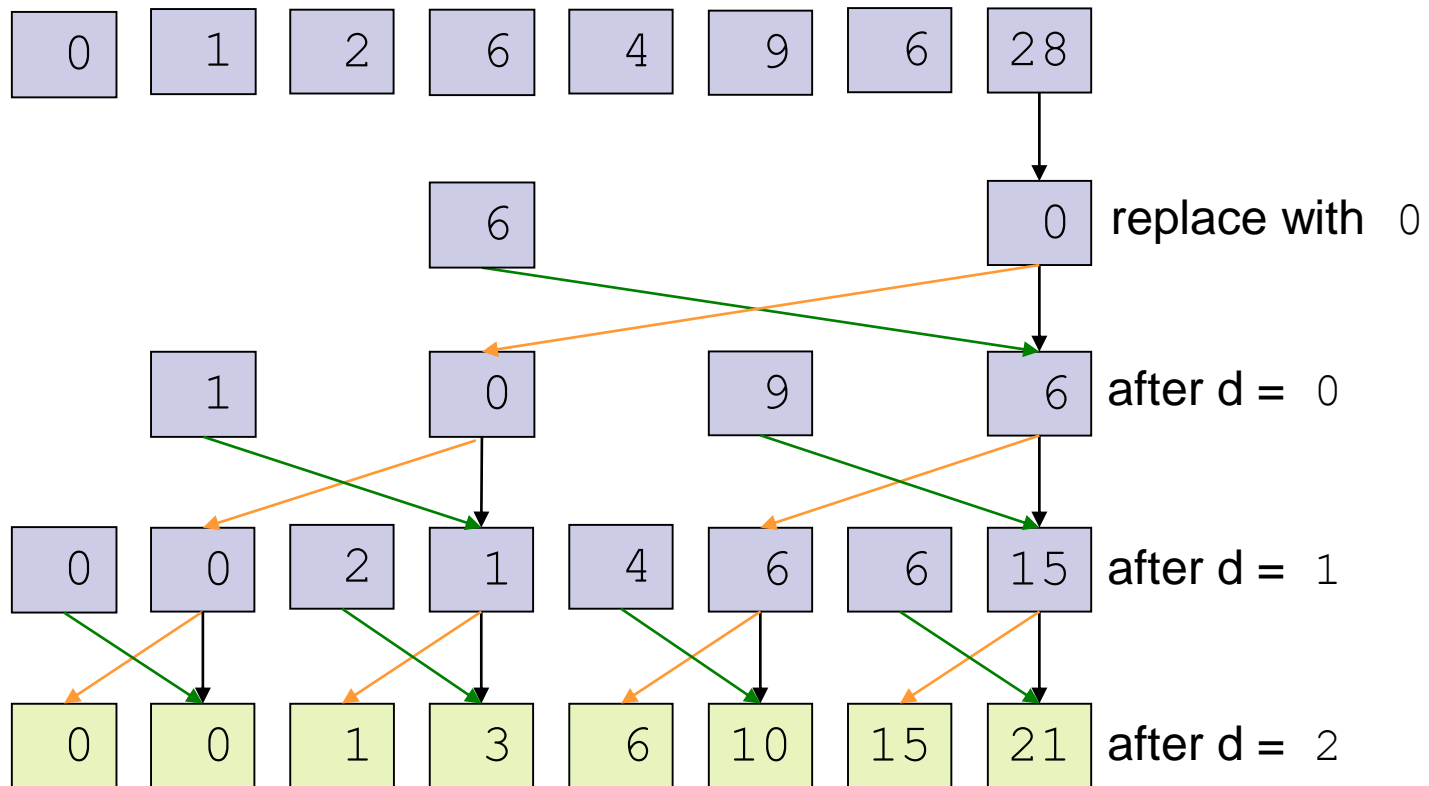
Work-Efficient Parallel Scan

■ Down-Sweep



Work-Efficient Parallel Scan

■ Down-Sweep





Work-Efficient Parallel Scan

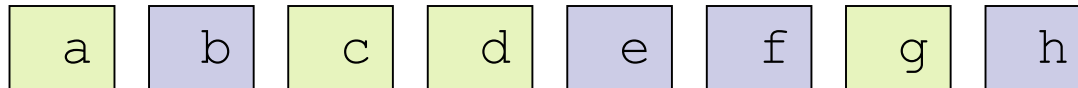
- Up-Sweep
 - $O(n)$ adds
- Down-Sweep
 - $O(n)$ adds
 - $O(n)$ swaps

Stream Compaction

■ *Stream Compaction*

□ Given an array of elements

- Create a new array with elements that meet a certain criteria, e.g. non null
- Preserve order

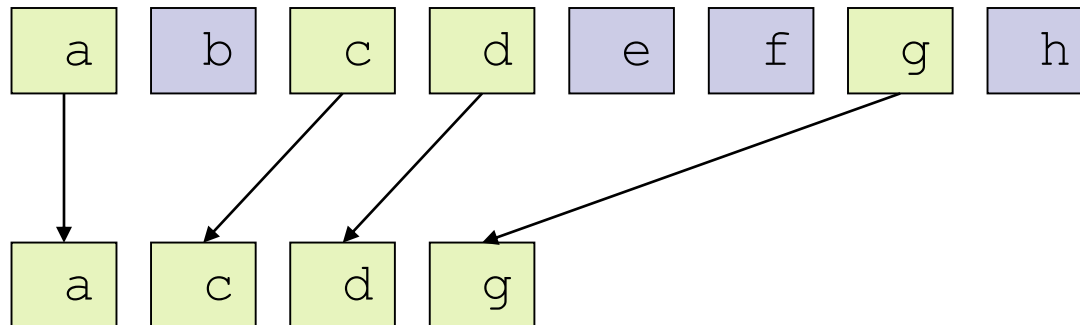


Stream Compaction

■ Stream Compaction

□ Given an array of elements

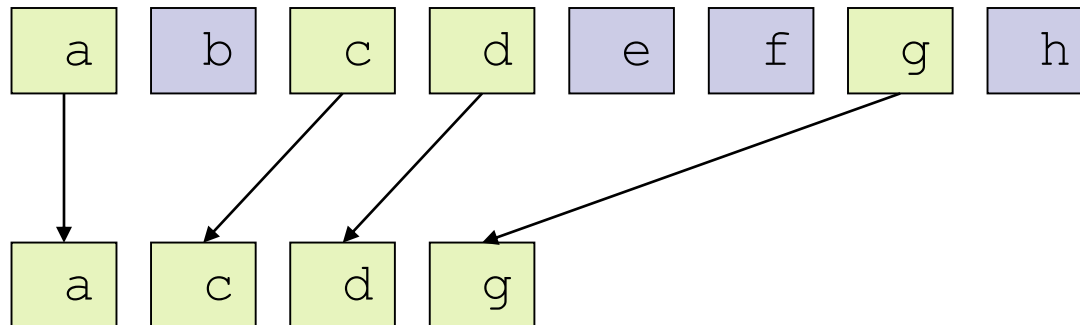
- Create a new array with elements that meet a certain criteria, e.g. non null
- Preserve order



Stream Compaction

■ Stream Compaction

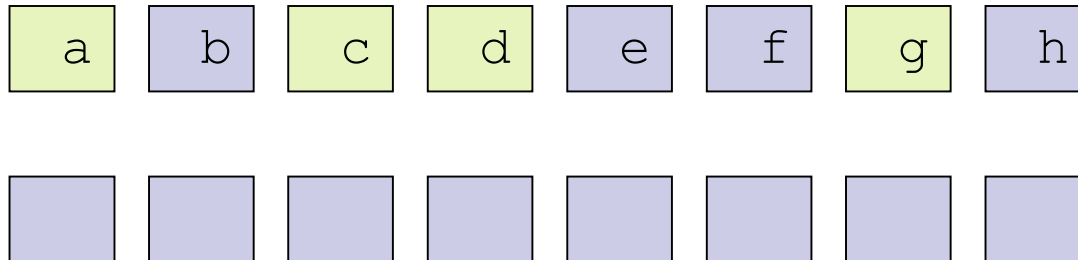
- Used in path tracing, collision detection, sparse matrix compression, etc.
- Can reduce bandwidth from GPU to CPU



Stream Compaction

■ Stream Compaction

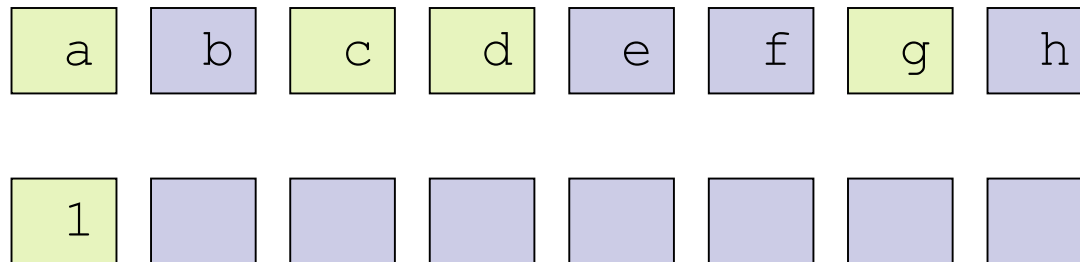
- *Step 1*: Compute temporary array containing
 - 1 if corresponding element meets criteria
 - 0 if element does not meet criteria



Stream Compaction

- Stream Compaction

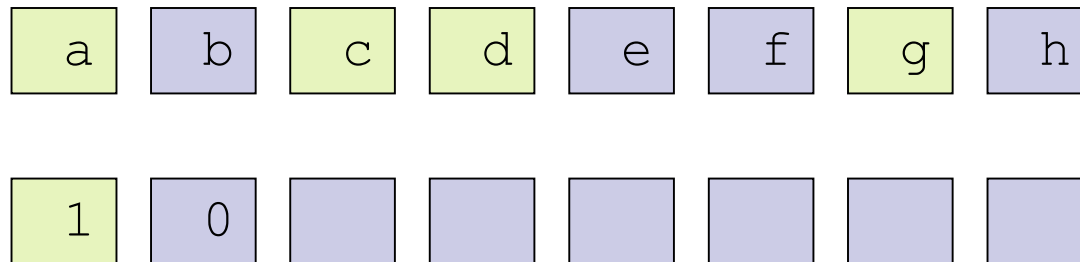
- *Step 1*: Compute temporary array



Stream Compaction

- Stream Compaction

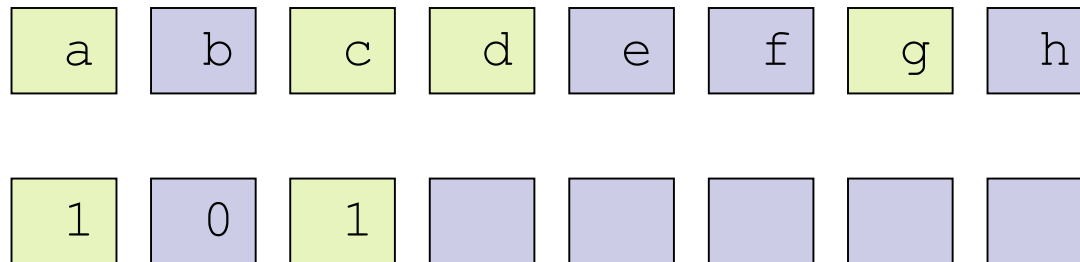
- *Step 1*: Compute temporary array



Stream Compaction

- Stream Compaction

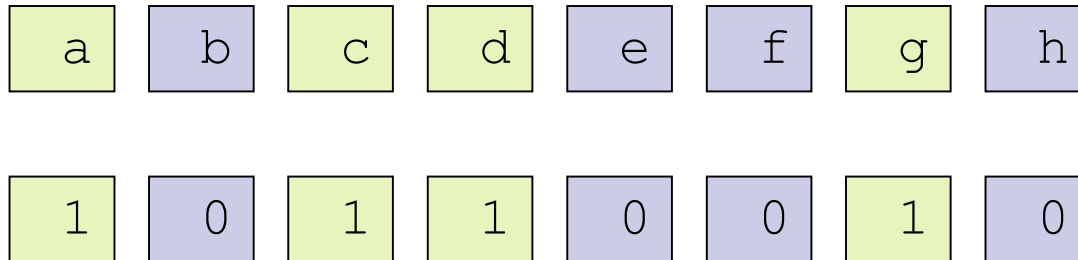
- *Step 1*: Compute temporary array



Stream Compaction

- Stream Compaction

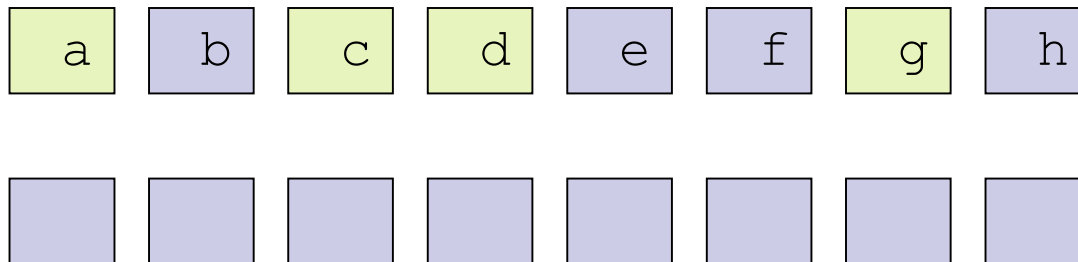
- *Step 1*: Compute temporary array



Stream Compaction

- Stream Compaction

- *Step 1*: Compute temporary array

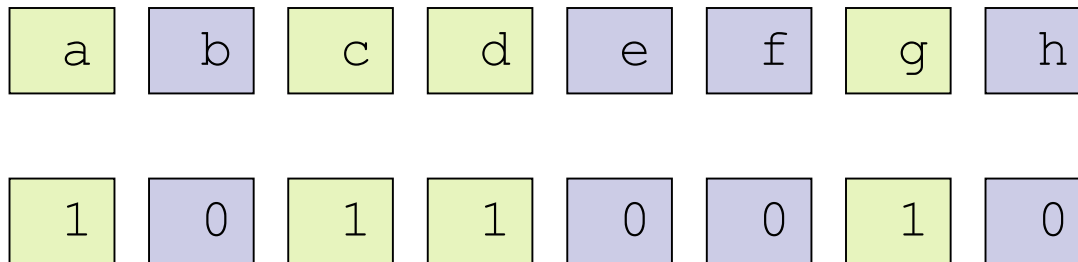


- It runs in parallel!

Stream Compaction

- Stream Compaction

- *Step 1*: Compute temporary array



- It runs in parallel!

Stream Compaction

■ Stream Compaction

- *Step 2*: Run exclusive scan on temporary array

	a	b	c	d	e	f	g	h
	1	0	1	1	0	0	1	0
Scan result:								

Stream Compaction

■ Stream Compaction

- *Step 2:* Run **exclusive** scan on temporary array

	a	b	c	d	e	f	g	h
	1	0	1	1	0	0	1	0
Scan result:	0	1	1	2	3	3	3	4

- Scan runs in parallel
- What can we do with the results?

Stream Compaction

■ Stream Compaction

□ *Step 3*: Scatter

- Result of scan is index into final array
- Only write an element if temporary array has a 1

Stream Compaction

■ Stream Compaction

□ *Step 3:* Scatter

	a	b	c	d	e	f	g	h
	1	0	1	1	0	0	1	0
Scan result:	0	1	1	2	3	3	3	4

Final array:				
	0	1	2	3

Stream Compaction

■ Stream Compaction

□ *Step 3:* Scatter

Scan result:

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
0	1	1	2	3	3	3	4

Final array:

a			
0	1	2	3

Stream Compaction

■ Stream Compaction

□ *Step 3:* Scatter

Scan result:

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
0	1	1	2	3	3	3	4

Final array:

a	c		
0	1	2	3

Stream Compaction

■ Stream Compaction

□ *Step 3:* Scatter

Scan result:

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
0	1	1	2	3	3	3	4

Final array:

a	c	d	
0	1	2	3

Stream Compaction

■ Stream Compaction

□ *Step 3:* Scatter

Scan result:

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
0	1	1	2	3	3	3	4

Final array:

a	c	d	g
0	1	2	3

Stream Compaction

■ Stream Compaction

□ *Step 3:* Scatter

	a	b	c	d	e	f	g	h
	1	0	1	1	0	0	1	0
Scan result:	0	1	1	2	3	3	3	4

Final array:				
	0	1	2	3

■ Scatter runs in parallel!

Stream Compaction

■ Stream Compaction

□ *Step 3:* Scatter

Scan result:

a	b	c	d	e	f	g	h
1	0	1	1	0	0	1	0
0	1	1	2	3	3	3	4

Final array:

a	c	d	g
0	1	2	3

■ Scatter runs in parallel!

Summed Area Table

- Summed *A*rea *T*able (*SAT*): 2D table where each element stores the sum of all elements in an input image between the lower left corner and the entry location.

Summed Area Table

■ Example:

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

4	9	12	14
2	6	9	11
2	5	6	8
1	2	2	4

$$(1 + 1 + 0) + (1 + 2 + 1) + (0 + 1 + 2) = 9$$

Summed Area Table

■ Benefit

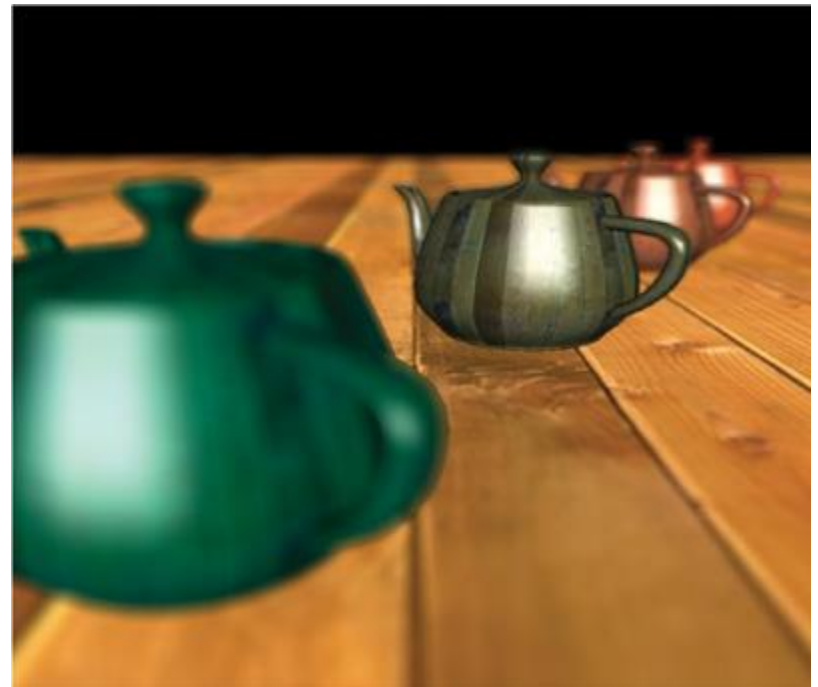
- Used to perform different width filters at every pixel in the image in constant time per pixel
- Just sample four pixels in SAT:

$$s_{filter} = \frac{s_{ur} - s_{ul} - s_{lr} + s_{ll}}{w \times h},$$

Summed Area Table

■ Uses

- Approximate depth of field
- Glossy environment reflections and refractions



Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

1			

Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

1	2		

Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

1	2	2	

Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

1	2	2	4

Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

2			
1	2	2	4

Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

2	5		
1	2	2	4



Summed Area Table

■ ■ ■

Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

4	9		
2	6	9	11
2	5	6	8
1	2	2	4

Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

4	9	12	
2	6	9	11
2	5	6	8
1	2	2	4

Summed Area Table

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

SAT

4	9	12	14
2	6	9	11
2	5	6	8
1	2	2	4



Summed Area Table

How would implement
this on the GPU?



Summed Area Table

How would compute a
SAT on the GPU using
inclusive scan?

Summed Area Table

■ Step 1 of 2:

Input image

2	1	0	0
0	1	2	0
1	2	1	0
1	1	0	2

Partial SAT

2	3	3	3
0	1	3	3
1	3	4	4
1	2	2	4



One inclusive scan for each row

Summed Area Table

■ Step 2 of 2:

Partial SAT

2	3	3	3
0	1	3	3
1	3	4	4
1	2	2	4

Final SAT

4	9	12	14
2	6	9	11
2	5	6	8
1	2	2	4



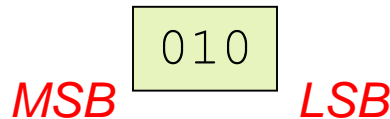
One inclusive scan for each column, bottom to top

Radix Sort

- Efficient for small sort keys
 - **k-bit** keys require k passes

Radix Sort

- Each radix sort pass partitions its input based on one bit
- First pass starts with the *least significant bit* (*LSB*). Subsequent passes move towards the *most significant bit* (*MSB*)



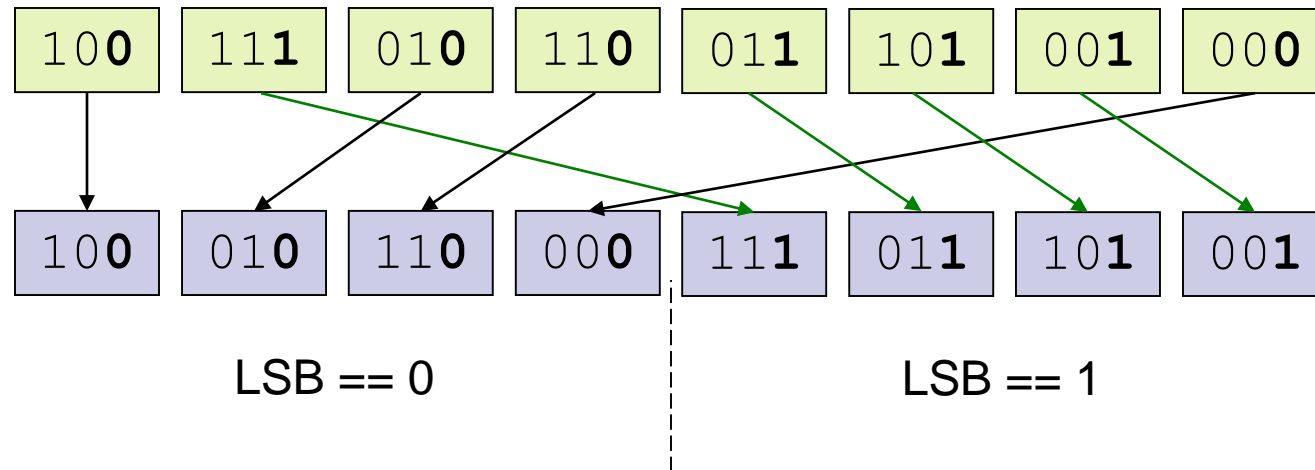
Radix Sort

- Example input:

100	111	010	110	011	101	001	000
-----	-----	-----	-----	-----	-----	-----	-----

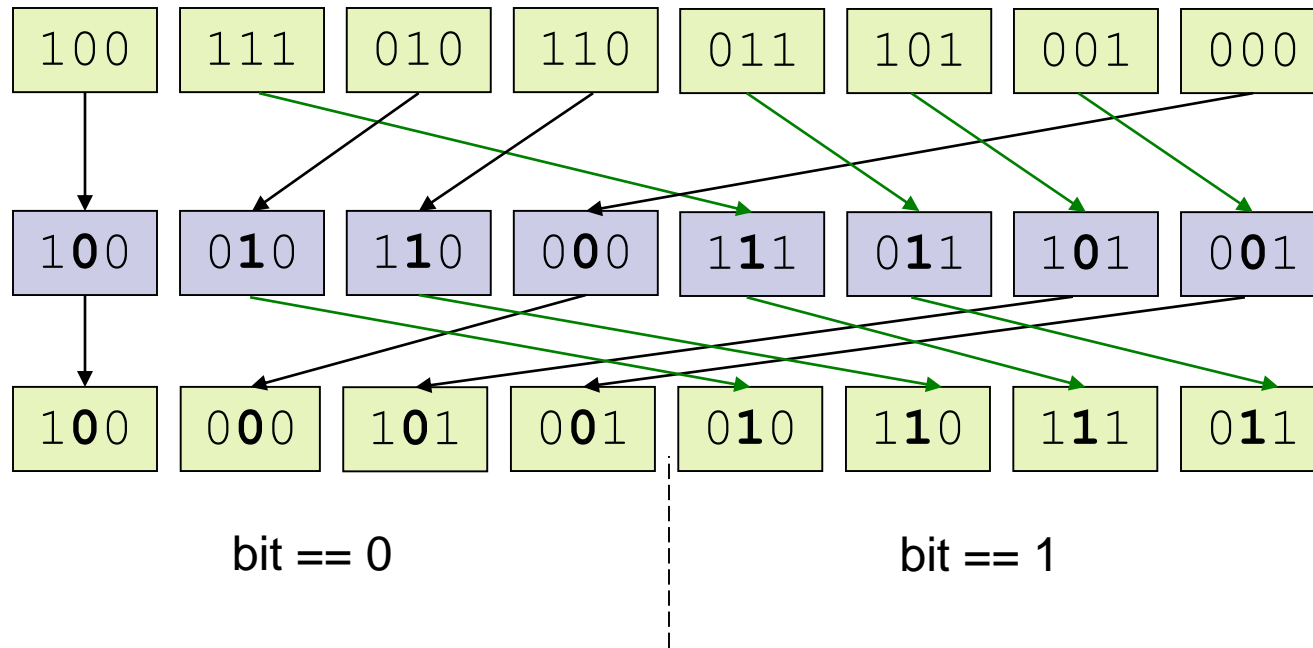
Radix Sort

- First pass: partition based on LSB



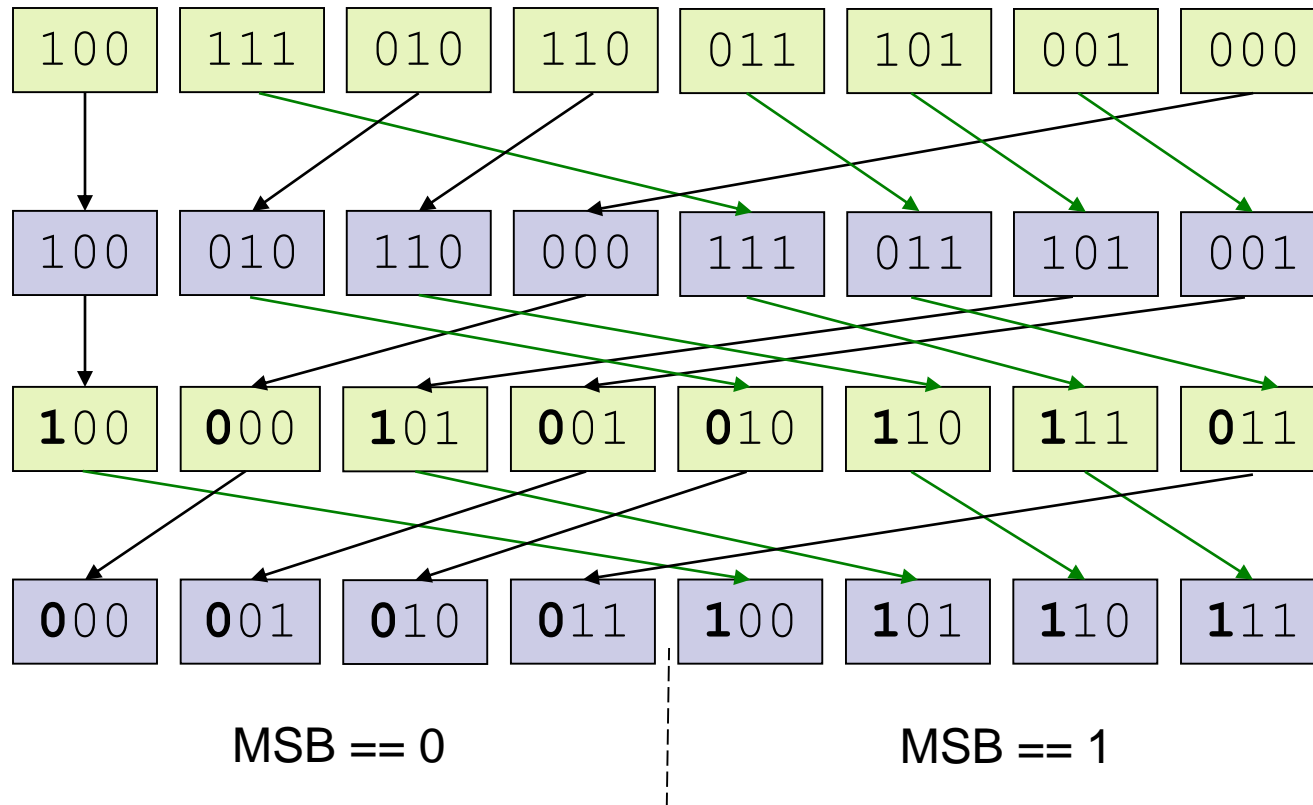
Radix Sort

- Second pass: partition based on middle bit



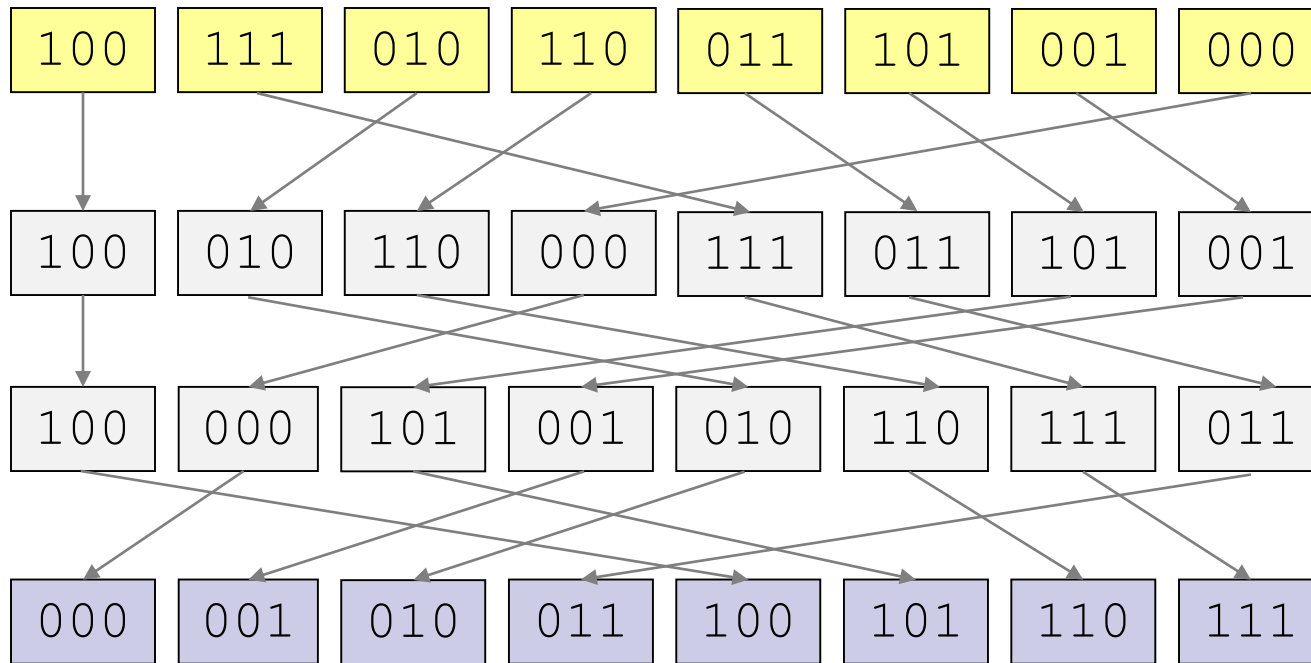
Radix Sort

- Final pass: partition based on MSB



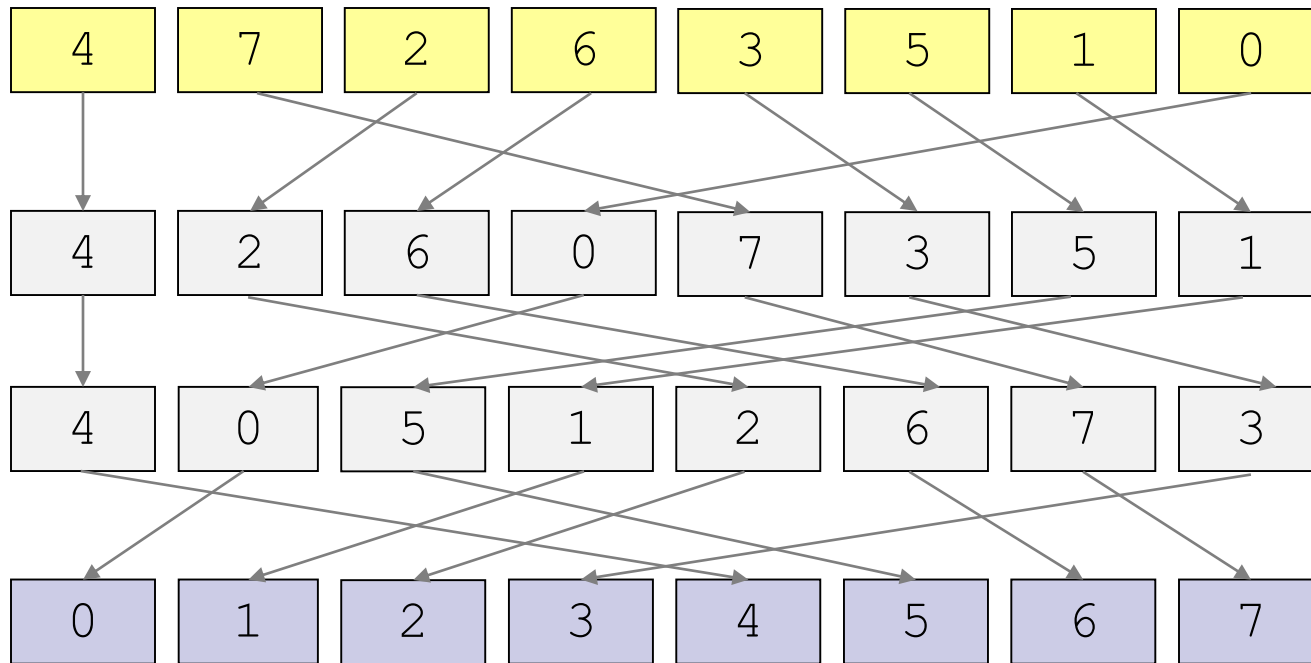
Radix Sort

■ Completed:



Radix Sort

■ Completed:





Parallel Radix Sort

- Where is the parallelism?

Parallel Radix Sort

1. Break input arrays into tiles
 - Each tile fits into shared memory for an SM
2. Sort tiles in *parallel* with *radix sort*
3. Merge pairs of tiles using a *parallel bitonic merge* until all tiles are merged.

Our focus is on Step 2



Parallel Radix Sort

- Where is the parallelism?
 - Each tile is sorted in parallel
 - Where is the parallelism within a tile?



Parallel Radix Sort

- Where is the parallelism?
 - Each tile is sorted in parallel
 - Where is the parallelism within a tile?
 - Each pass is done in sequence after the previous pass. No parallelism
 - Can we parallelize an individual pass? How?
 - Merge also has parallelism

Parallel Radix Sort

- Implement *spilt*. Given:

- Array, *i*, at pass *n*:

100	111	010	110	011	101	001	000
-----	-----	-----	-----	-----	-----	-----	-----

- Array, *b*, which is true/false for bit *n*:

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

- Output array with false keys before true keys:

100	010	110	000	111	011	101	001
-----	-----	-----	-----	-----	-----	-----	-----

Parallel Radix Sort

■ Step 1: Compute *e* array

100	111	010	110	011	101	001	000	i array
0	1	0	0	1	1	1	0	b array
<hr/>								
1	0	1	1	0	0	0	1	e array

Parallel Radix Sort

■ Step 2: Exclusive Scan *e*

100	111	010	110	011	101	001	000	i array
0	1	0	0	1	1	1	0	b array
<hr/>								
1	0	1	1	0	0	0	1	e array
0	1	1	2	3	3	3	3	f array

Parallel Radix Sort

■ Step 3: Compute *totalFalses*

100	111	010	110	011	101	001	000	i array
0	1	0	0	1	1	1	0	b array
<hr/>								
1	0	1	1	0	0	0	1	e array
0	1	1	2	3	3	3	3	f array

Chuyển bit 0 có sign = 1
suy ra lúc scan lại thì sẽ biết phía trước
bao nhiêu bit 0

$$\begin{aligned}\text{totalFalses} &= e[n-1] + f[n-1] \\ \text{totalFalses} &= 1 + 3 \\ \text{totalFalses} &= 4\end{aligned}$$

Parallel Radix Sort

■ Step 4: Compute *t* array

100	111	010	110	011	101	001	000	i array
0	1	0	0	1	1	1	0	b array
<hr/>								
1	0	1	1	0	0	0	1	e array
0	1	1	2	3	3	3	3	f array
								t array

$$t[i] = i - f[i] + \text{totalFalses}$$

totalFalses = 4

Parallel Radix Sort

■ Step 4: Compute *t* array

100	111	010	110	011	101	001	000	i array
0	1	0	0	1	1	1	0	b array
<hr/>								
1	0	1	1	0	0	0	1	e array
0	1	1	2	3	3	3	3	f array
4								t array

$$t[0] = 0 - f[0] + \text{totalFalses}$$

$$t[0] = 0 - 0 + 4$$

$$t[0] = 4$$

0: i

f[0]: số thàng 0 phía trước

4: tổng số thàng 0

totalFalses = 4

Parallel Radix Sort

■ Step 4: Compute *t* array

100	111	010	110	011	101	001	000	i array
0	1	0	0	1	1	1	0	b array
<hr/>								
1	0	1	1	0	0	0	1	e array
0	1	1	2	3	3	3	3	f array
4	4							t array

$$t[1] = 1 - f[1] + \text{totalFalses}$$

$$t[1] = 1 - 1 + 4$$

$$t[1] = 4$$

$$\text{totalFalses} = 4$$

Parallel Radix Sort

■ Step 4: Compute *t* array

100	111	010	110	011	101	001	000	i array
0	1	0	0	1	1	1	0	b array
<hr/>								
1	0	1	1	0	0	0	1	e array
0	1	1	2	3	3	3	3	f array
4	4	5						t array

$$t[2] = 2 - f[2] + \text{totalFalses}$$

$$t[2] = 2 - 1 + 4$$

$$t[2] = 5$$

$$\text{totalFalses} = 4$$

Parallel Radix Sort

■ Step 4: Compute *t* array

100	111	010	110	011	101	001	000	i array
0	1	0	0	1	1	1	0	b array
<hr/>								
1	0	1	1	0	0	0	1	e array
0	1	1	2	3	3	3	3	f array
4	4	5	5	5	6	7	8	t array

$$t[i] = i - f[i] + \text{totalFalses}$$

$$\text{totalFalses} = 4$$

Parallel Radix Sort

■ Step 5: Scatter based on address *d*

100	111	010	110	011	101	001	000
-----	-----	-----	-----	-----	-----	-----	-----

 i array

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

 b array

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

 e array

0	1	1	2	3	3	3	3
---	---	---	---	---	---	---	---

 f array

4	4	5	5	5	6	7	8
---	---	---	---	---	---	---	---

 t array

0							
---	--	--	--	--	--	--	--

 $d[i] = b[i] ? t[i] : f[i]$

Parallel Radix Sort

■ Step 5: Scatter based on address *d*

100	111	010	110	011	101	001	000
-----	-----	-----	-----	-----	-----	-----	-----

 i array

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

 b array

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

 e array

0	1	1	2	3	3	3	3
---	---	---	---	---	---	---	---

 f array

4	4	5	5	5	6	7	8
---	---	---	---	---	---	---	---

 t array

0	4						
---	---	--	--	--	--	--	--

 $d[i] = b[i] ? t[i] : f[i]$

Parallel Radix Sort

■ Step 5: Scatter based on address *d*

100	111	010	110	011	101	001	000
-----	-----	-----	-----	-----	-----	-----	-----

 i array

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

 b array

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

 e array

0	1	1	2	3	3	3	3
---	---	---	---	---	---	---	---

 f array

4	4	5	5	5	6	7	8
---	---	---	---	---	---	---	---

 t array

0	4	1					
---	---	---	--	--	--	--	--

 $d[i] = b[i] ? t[i] : f[i]$

Parallel Radix Sort

■ Step 5: Scatter based on address *d*

100	111	010	110	011	101	001	000
-----	-----	-----	-----	-----	-----	-----	-----

 i array

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

 b array

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

 e array

0	1	1	2	3	3	3	3
---	---	---	---	---	---	---	---

 f array

4	4	5	5	5	6	7	8
---	---	---	---	---	---	---	---

 t array

0	4	1	2	5	6	7	3
---	---	---	---	---	---	---	---

 $d[i] = b[i] ? t[i] : f[i]$

Parallel Radix Sort

■ Step 5: Scatter based on address *d*

100	111	010	110	011	101	001	000
-----	-----	-----	-----	-----	-----	-----	-----

 i array

0	4	1	2	5	6	7	3
---	---	---	---	---	---	---	---

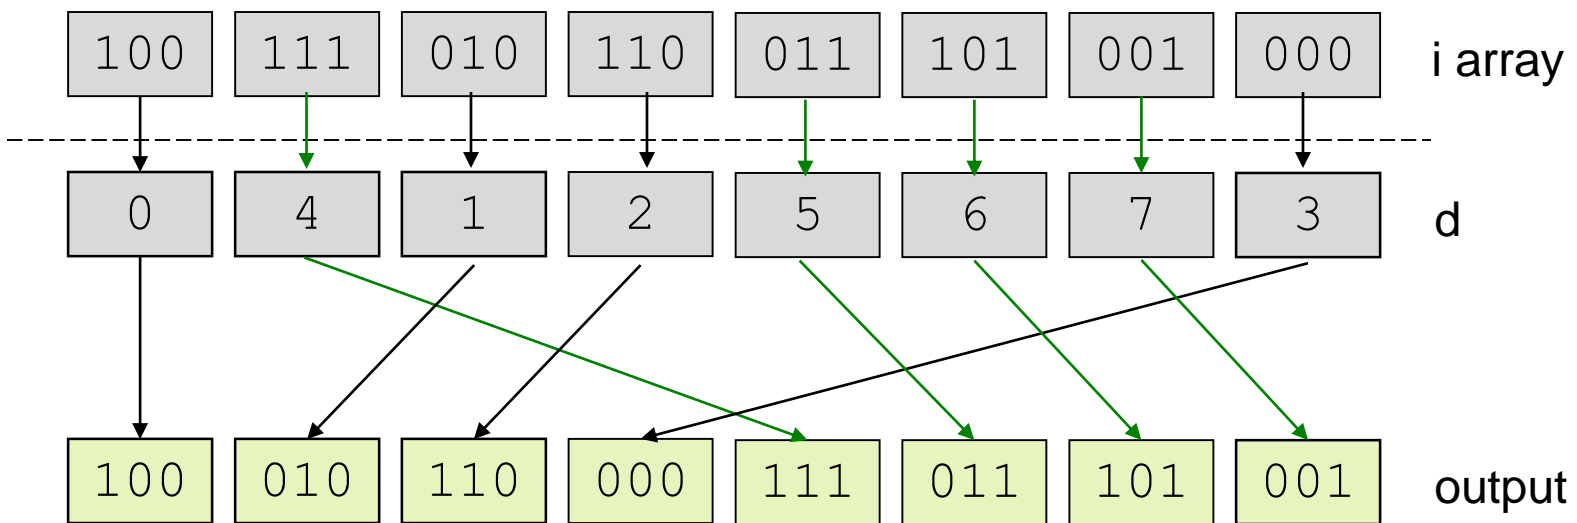
 d

--	--	--	--	--	--	--	--

 output

Parallel Radix Sort

■ Step 5: Scatter based on address *d*





Parallel Radix Sort

- Given k-bit keys, how do we sort using our new *split* function?
- Once each tile is sorted, how do we merge tiles to provide the final sorted array?



Summary

- Parallel reduction, scan, and sort are building blocks for many algorithms
- An understanding of parallel programming and GPU architecture yields efficient GPU implementations